# MLOps

Introduction

## Outline

- Implementation and Automation
  - Continuous Integration (CI)
  - Continuous Delivery (CD)
  - Continuous Monitoring & Training (CT)

Den Ope [ ciled] + em/ct => mrops

7. Integralin to Cross
Playfron. Design, Steaming. 2. Coding I sall stronge 3. Dering H. Pellar Dellar bymont

Prof. Sashikumaar Ganesan



#### Data Science & ML

- Capable of solving complex real-world problems, transforming businesses, and delivering values in all domains
- In forefront due to the access to
  - Large datasets
  - Inexpensive on-demand compute resources, cloud solutions
  - Advanced easy to use ML tools
  - Rapid advances in ML methods and applications
  - Many businesses investing in developing their own AI team

### Recent advances in Data Science & ML

- Demand for Data Engineers, Data Scientists and ML Engineers
- Emergence of MLOps
  - Unifies Al system developments and Al system operations
  - Automates and monitors all steps in AI system construction, integration, testing, releasing, deployment, data collection, cleaning, feature engineering, feature selection, retraining and infrastructure management, etc.

MLOPS :- Montre lan Opens

# Challenges

Al deeply depend on Cloud Computing

• Raw ingredients of AI requires massive compute, extensive data and specialized hardware

H 1100 - 500/h 200/h

# AI/ML Systems

- Al systems mainly focus on data engineering, data processing, problem feasibility and business alignment
- Primary focus is on business with ML rather than code "Highest Paid Person's Opinion" (HIPPO) effect
- Most Al systems are not code native, use academic software packages that do not scale for large-scale problems

INN MAG

# What is MLOps?

- The process of automating AI system using DevOps methodologies
- Share

#### Hidden Technical Debt in Machine Learning Systems

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips {dsculley, gholt, dgg, edavydov, toddphillips}@google.com Google, Inc.

Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, Dan Dennison {ebner, vchaudhary, mwyoung, jfcrespo, dennison}@google.com@Google, Inc.

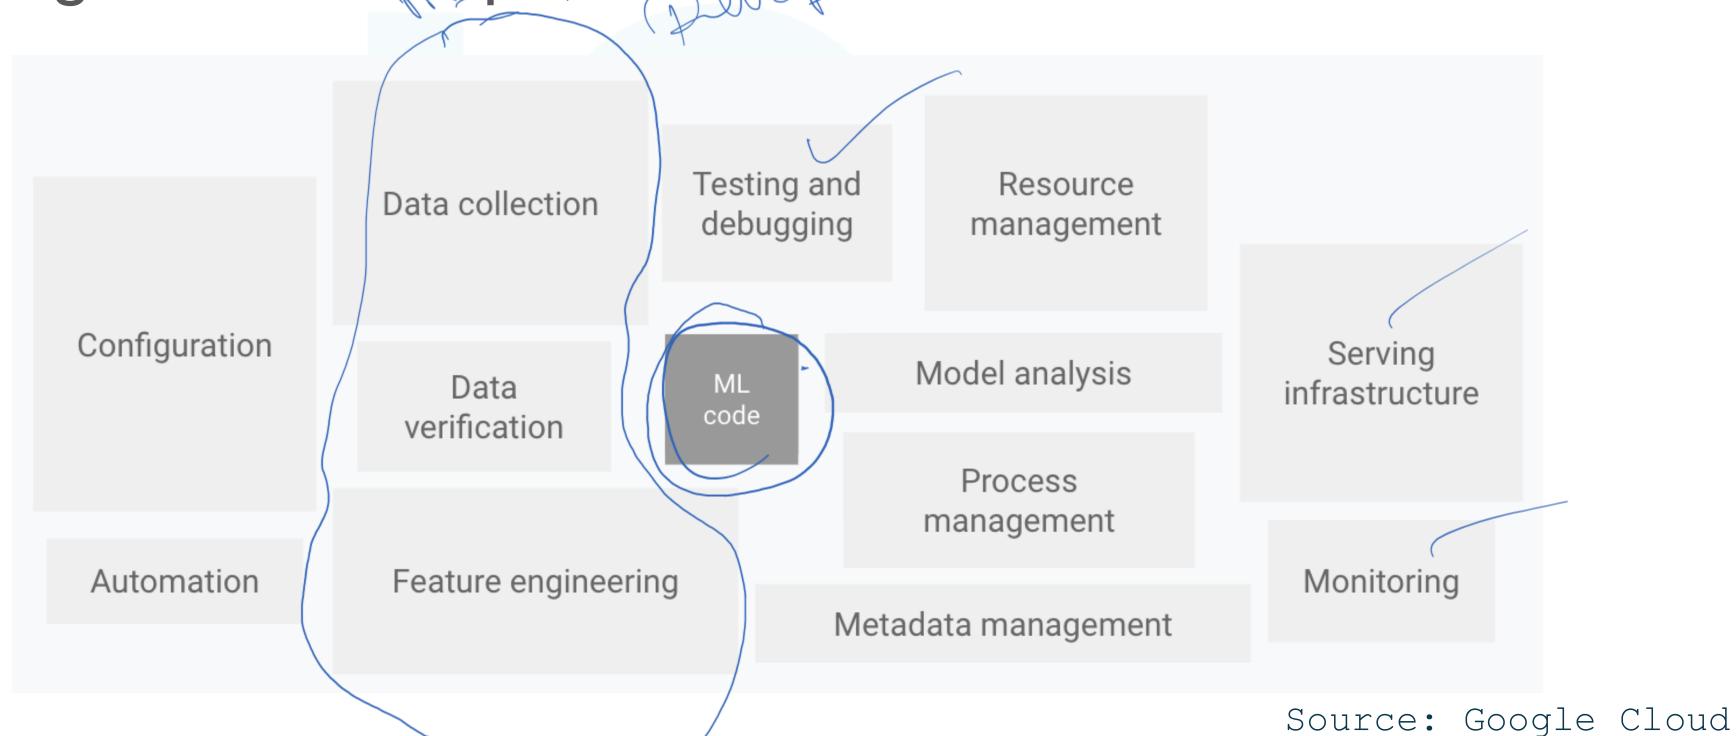
Source: Google Cloud

# What is MLOps?

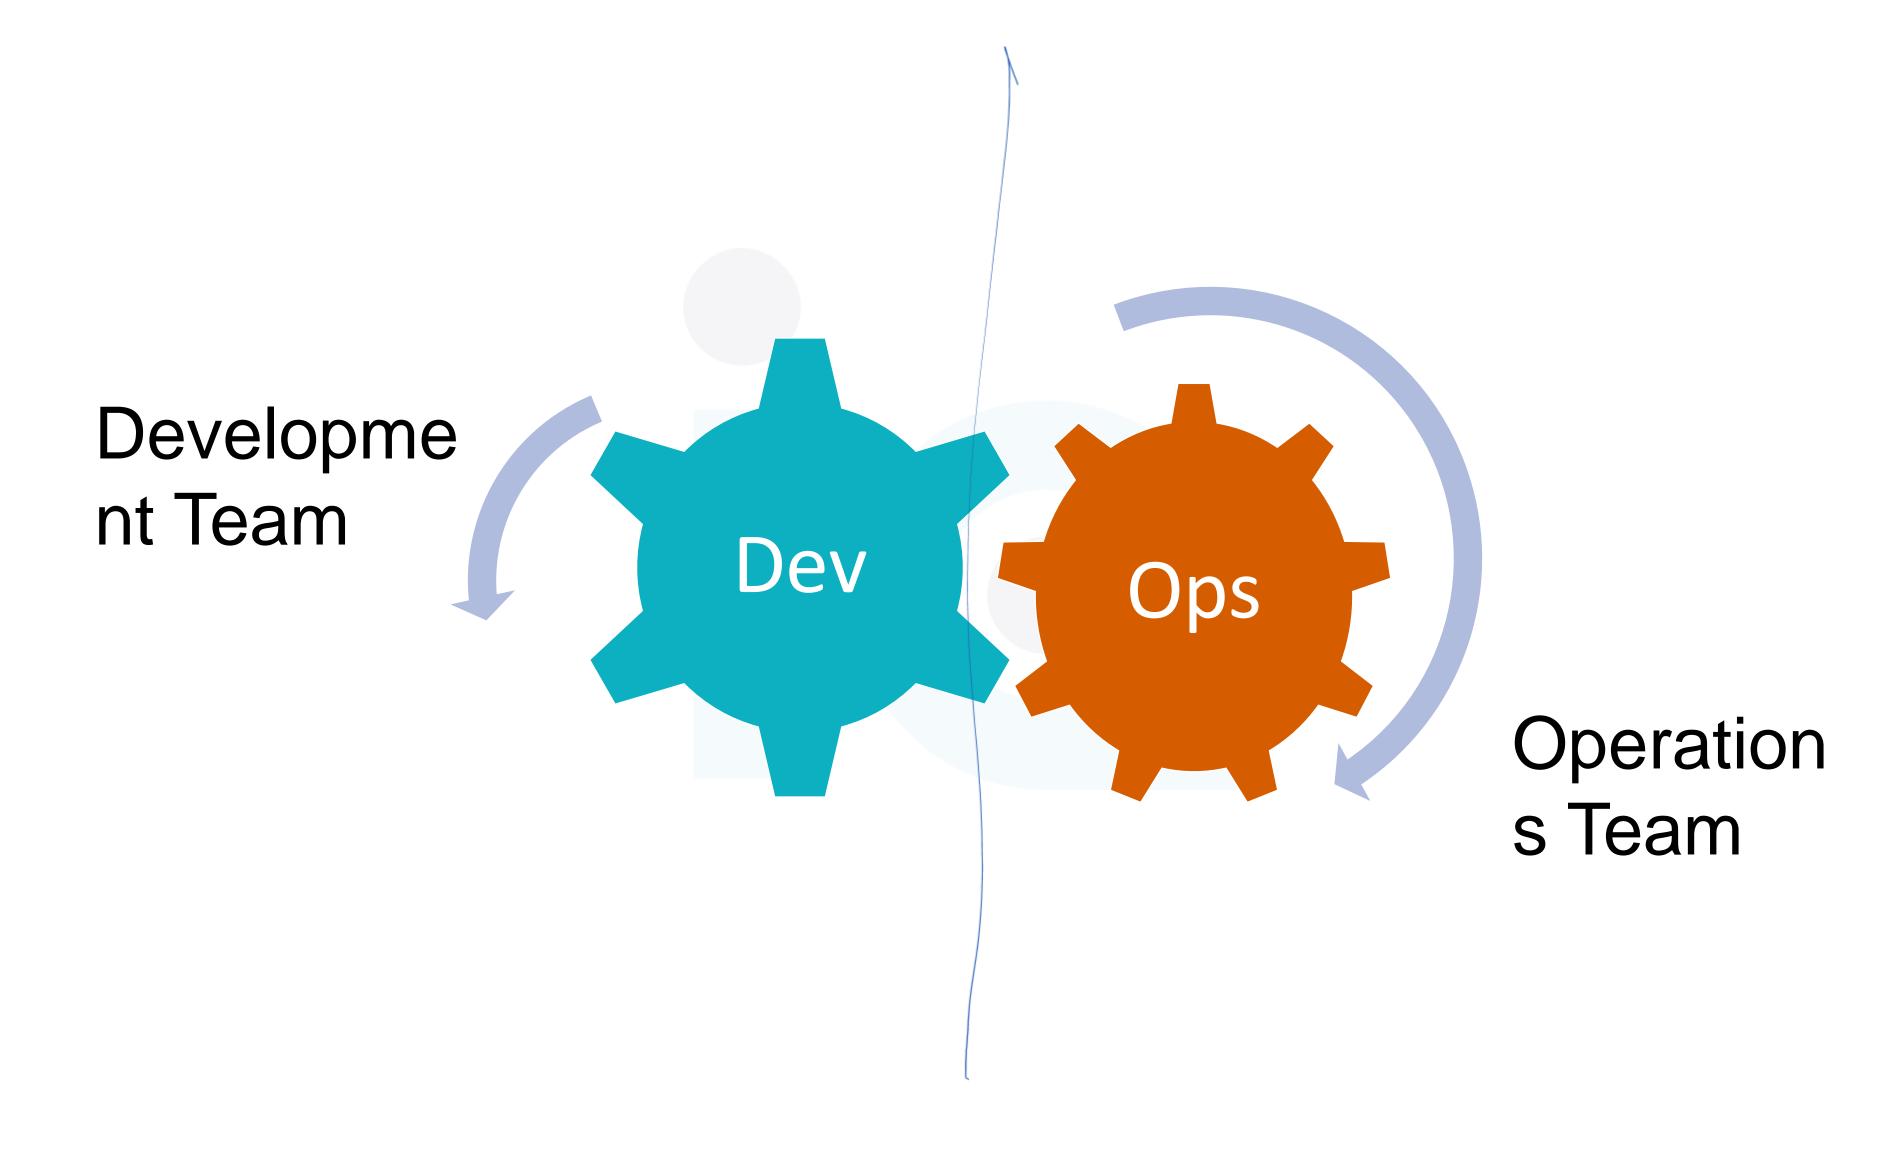
The process of automating ML systems using DevOps methodologies

• Shares lineage with DevOps, which demands automation

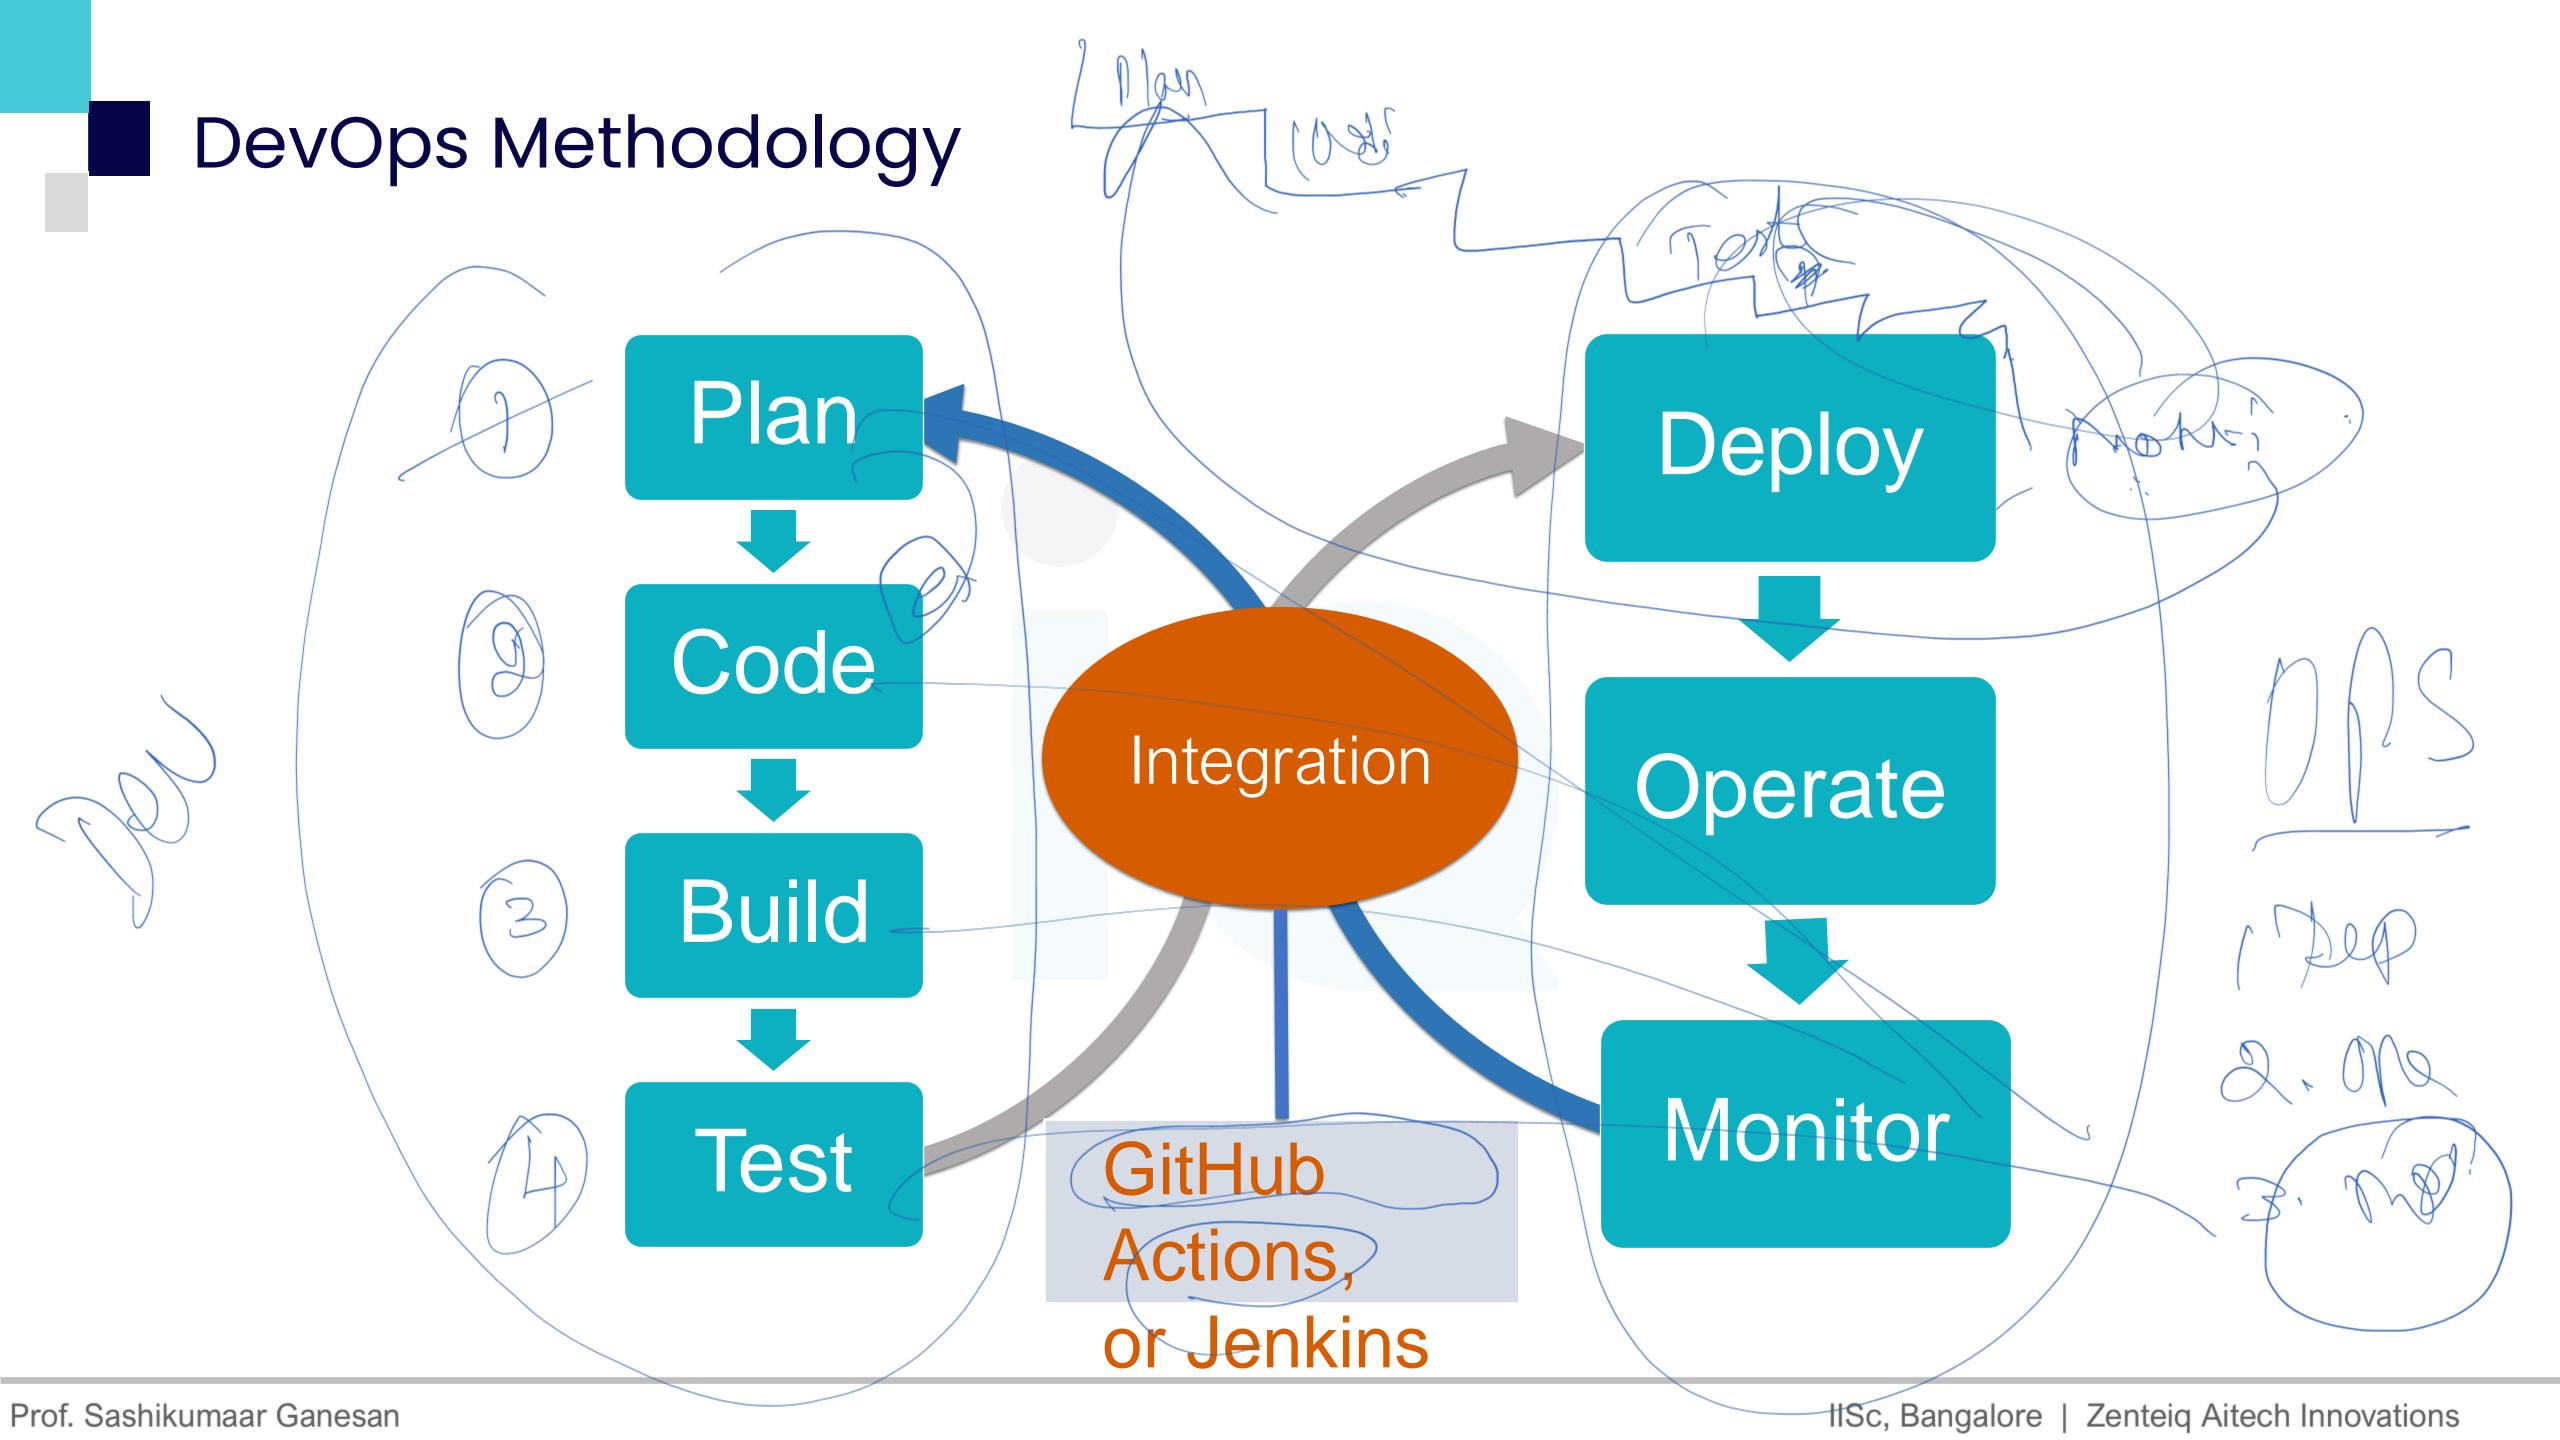
Elements of MLOps

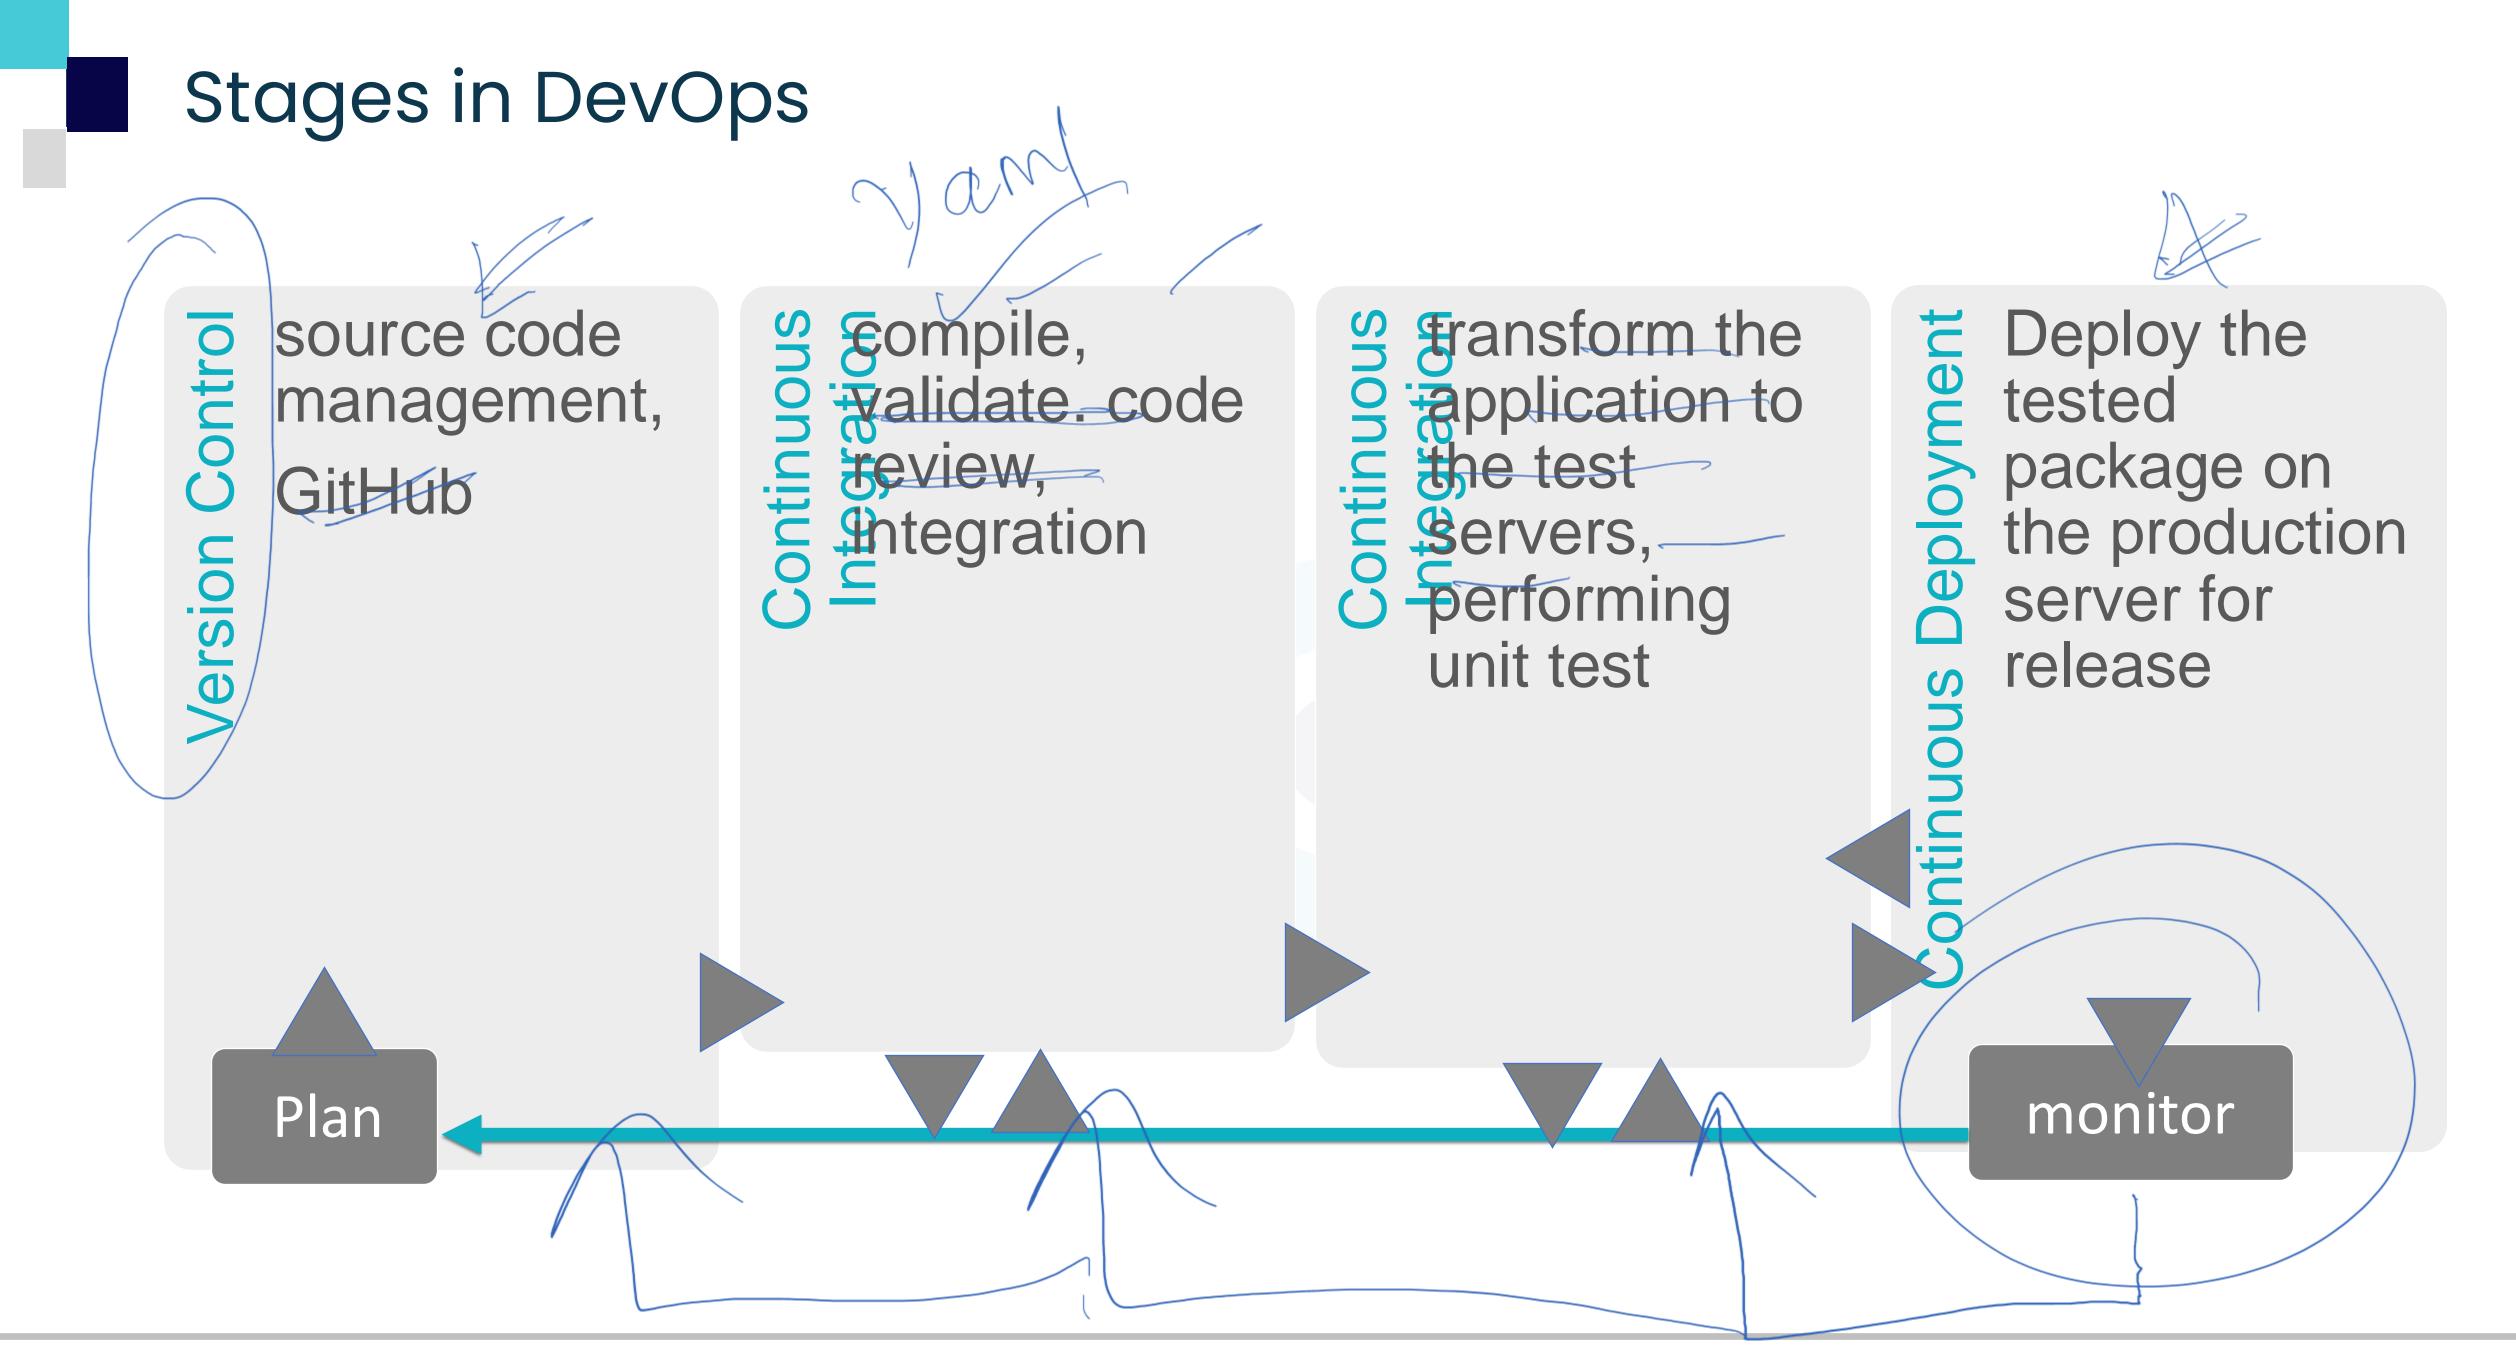


# DevOps Methodology



Prof. Sashikumaar Ganesan





# MLOps versus DevOps

- With MLOps, not only the software engineering process needs full automation, but so do the data and ML models
- Continuous Training
  - Model training and deployment is added to DevOps
- Monitor new things that break automation
  - Data-drift the delta changes between the data from last time the

model training occurred.

The property of the

# MLOps differs from DevOps!





Data
scientists:
Focusing
EDA, model
development,
&
experimentati
o, no
necessarily

software



# Developm ent

Experiment different features, algorithms, models, hyperparamet er selection, reproducibility



# **Testing**

Not software but ML system. In addition to unit and integration tests, data validation, trained model quality, ad model validation



## Deployme nt

Involves multi-step pipeline to automate retraining and deploy model



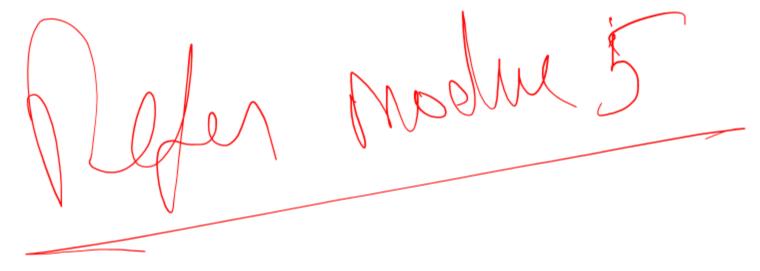
# Productio n

model decays in many ways than conventional software

# MLOps differ from software systems!

- CI is not only just testing and validation of software, but also data, data schemas and models
- CD is not only for single software or service but a system including deploying another system
- CM/CT is a new methodology, unique to MLOps, that retrains the ML model

# Eight Steps in AI/ML system Development



#### 1. Data Extraction

Select and integrate relevant data from various data source

### 2. Data Analysis

 Perform Exploratory data analysis (EDA) to understand the data

# 3. Data Preparation

 Prepare data (cleaning, split into training, validation & testing) for ML task

## 4. ML Model training

 Train appropriate ML model, tune hyperparameters



#### Overview

- Principles, best practices, and tools for scalable DataOps pipelines in ML.
- Importance: Handling large data volumes, real-time processing, resource optimization, system reliability.
- Data Ingestion: Tools Apache Kafka, Apache NiFi, Amazon Kinesis.
- Data Processing: ETL vs. ELT, Tools Apache Spark, Dask.
- Data Storage: Solutions HDFS, S3, Google Cloud Storage, Delta Lake.
- Data Quality: Tools Great Expectations, Apache Griffin.
- Architecture: Decoupling, modularity, fault tolerance.
- Orchestration: Tools Apache Airflow, Prefect, Apache NiFi.
- Monitoring: Tools Prometheus, Grafana, ELK stack.
- Case Study: Real-world scalable DataOps pipeline example.

Prof. Sashikumaar Ganesan | Zenteiq 2024 | MLOps



# Eight Steps in AI/ML system Development

- 5. Model evaluation
  - Evaluate model on "Holdout test set" and assess the model
- 6. Model validation
  - Is the trained model adequate for deployment?
  - Predictive performance is better than a baseline?
- 7. Al system serving
  - Deployment to a target environment (Microservices, Edge, etc)
- 8. Al system monitoring
  - Monitor for new iteration/retraining

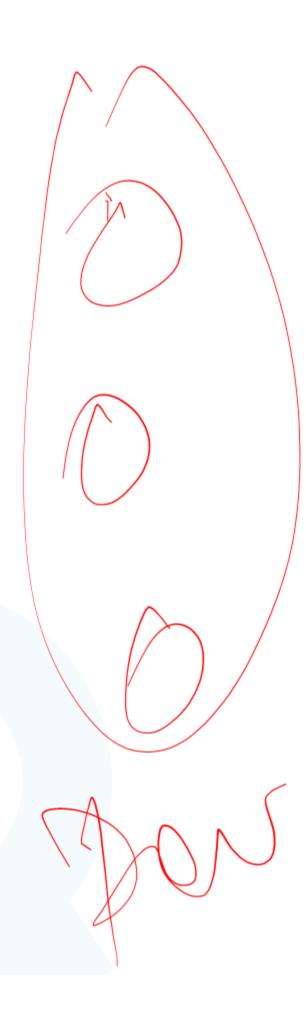
# MLOps: How to start?

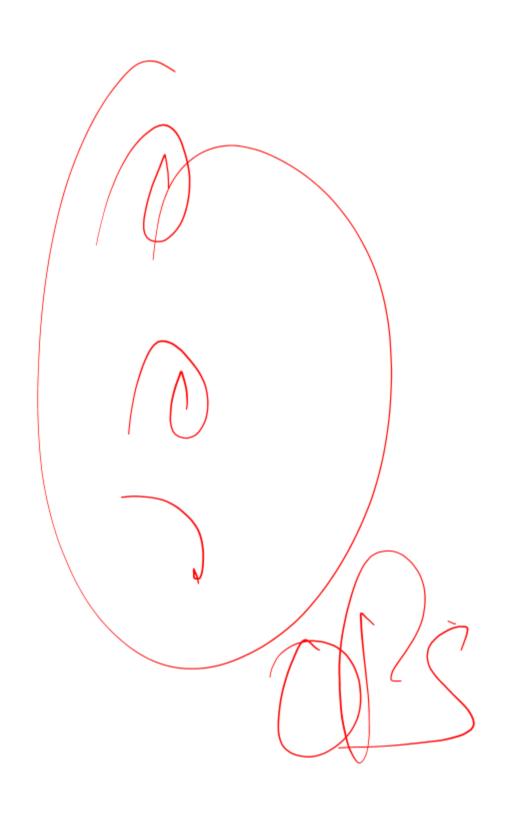


Level 1: ML Pipeline automation

Level 2: CI, CD and CT Pipeline

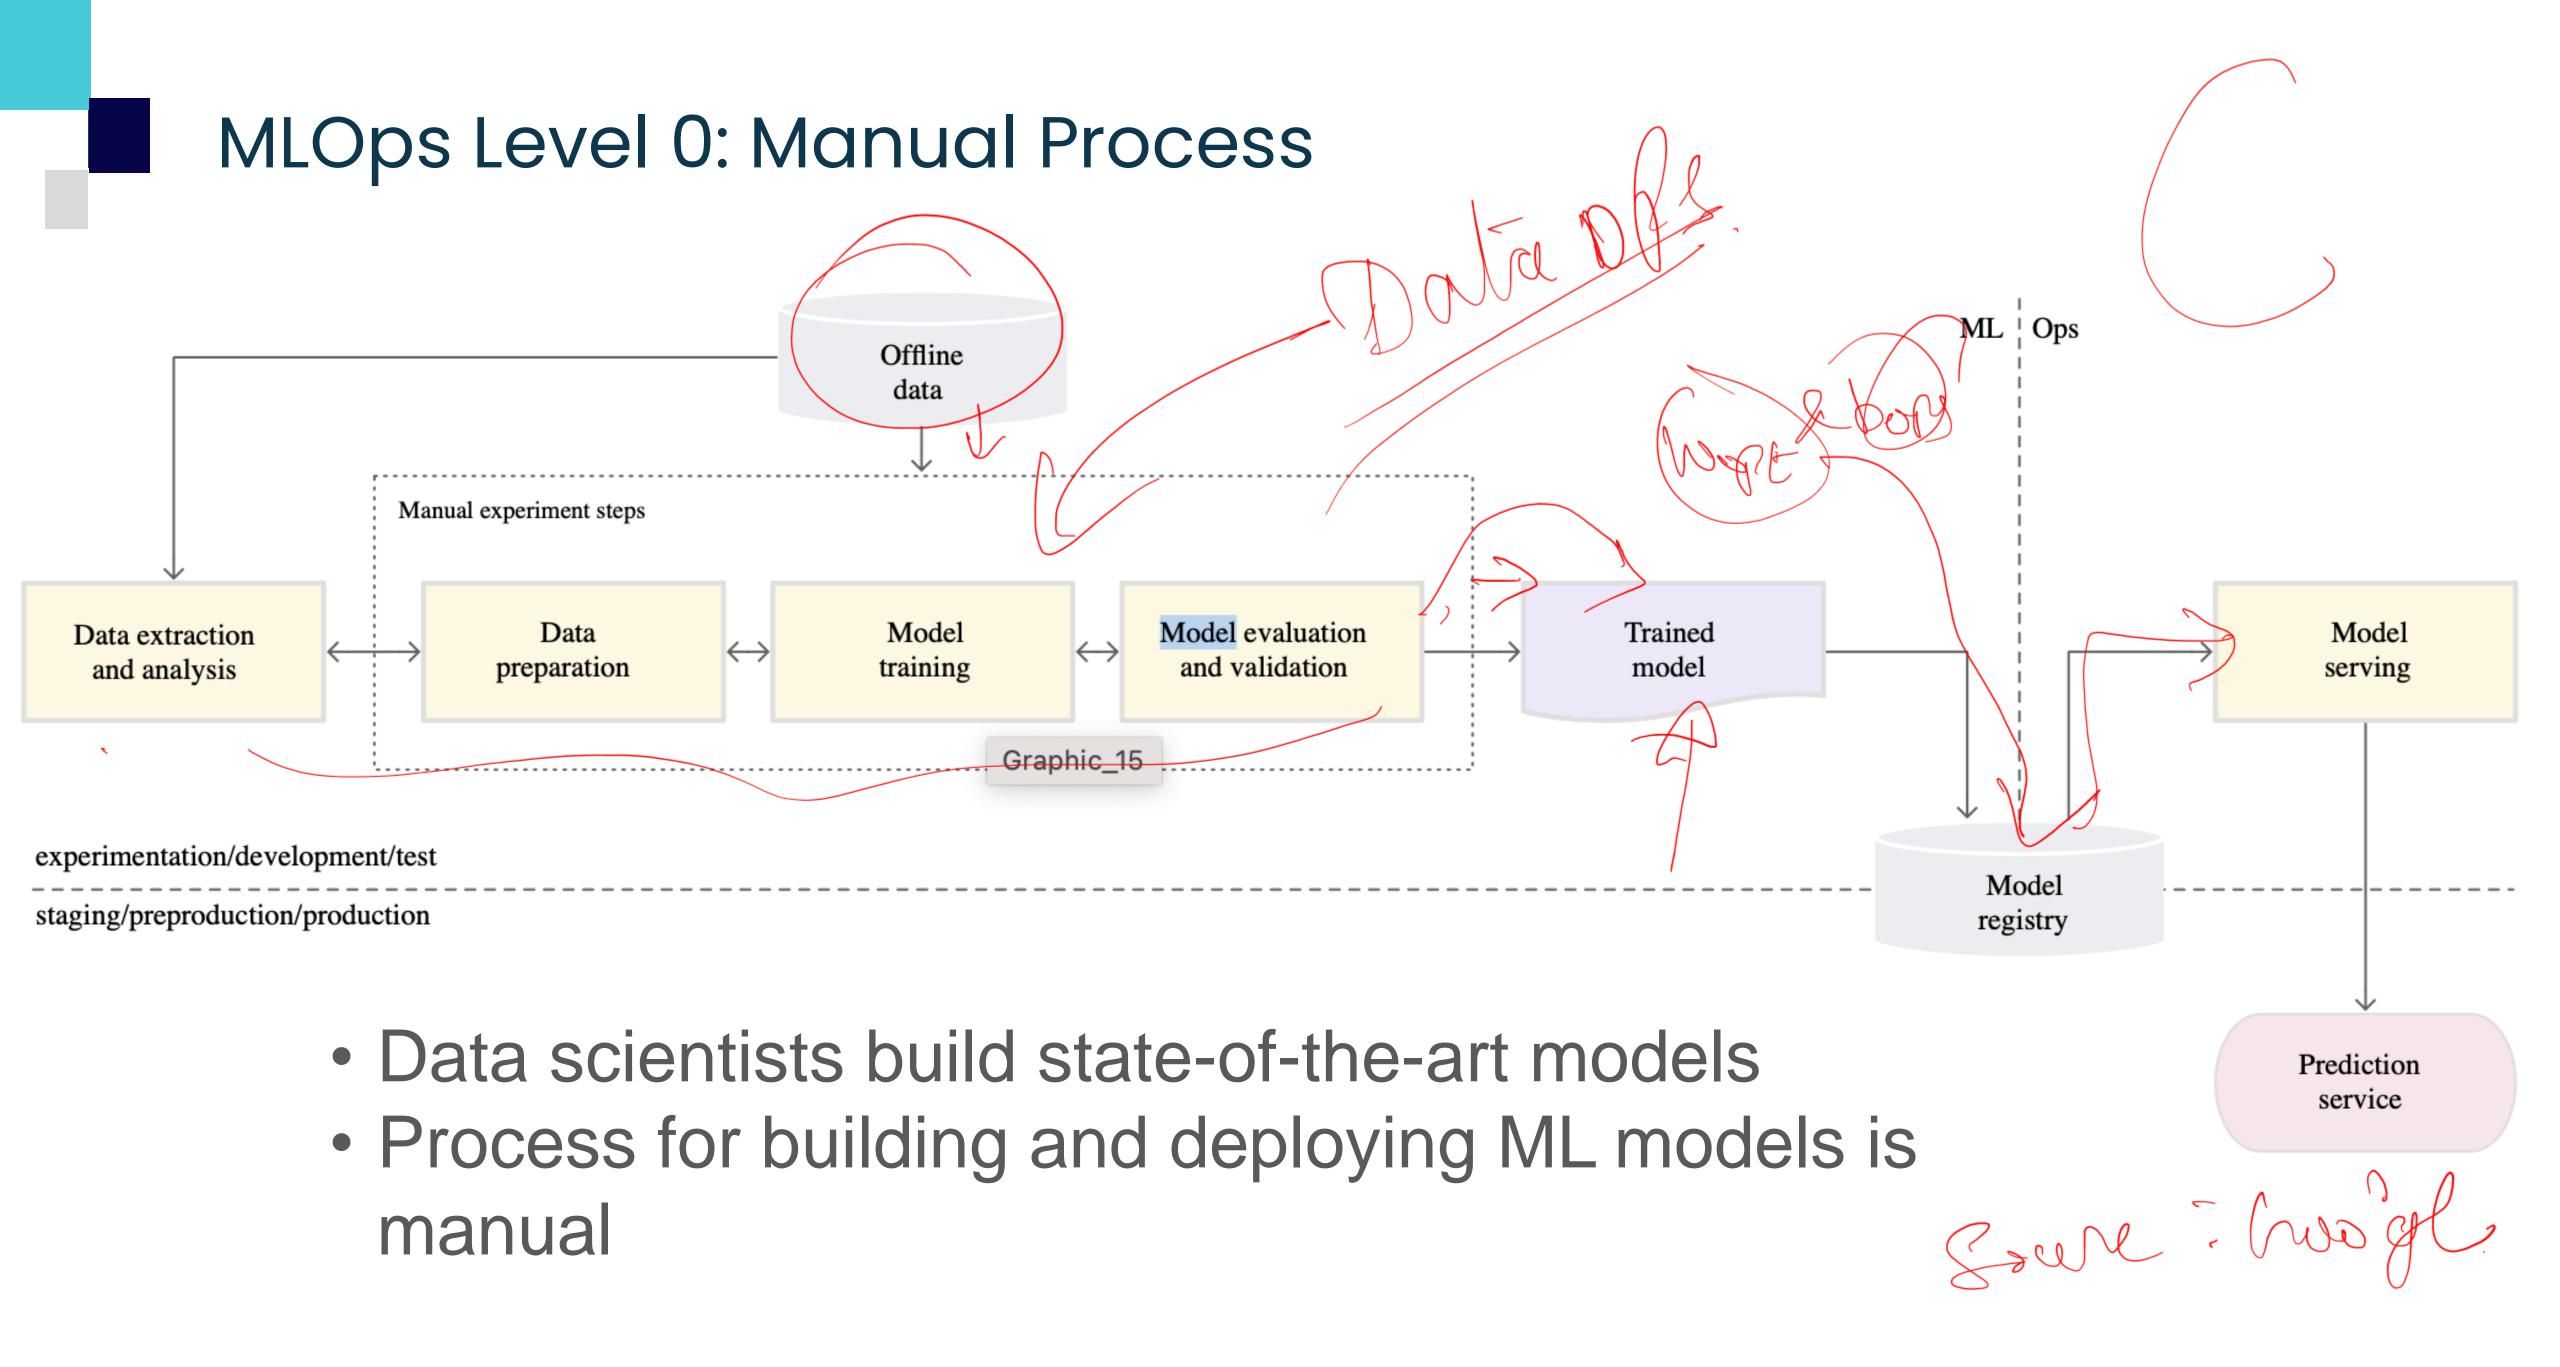
automation





# MLOps Level 0: Manual Process

- Manual, script-driven and iterative process
  - Every step is manual:
    - Data analysis, data preparation, model training, and validation
- Disconnect between ML and operations
  - Separates Data Engineers, Data scientists and ML Engineers
- Infrequent release or update or retraining
- No CI
- No CD
- Prediction service, not the full AI System, is referred to as deployment
- Lack of active performance monitoring



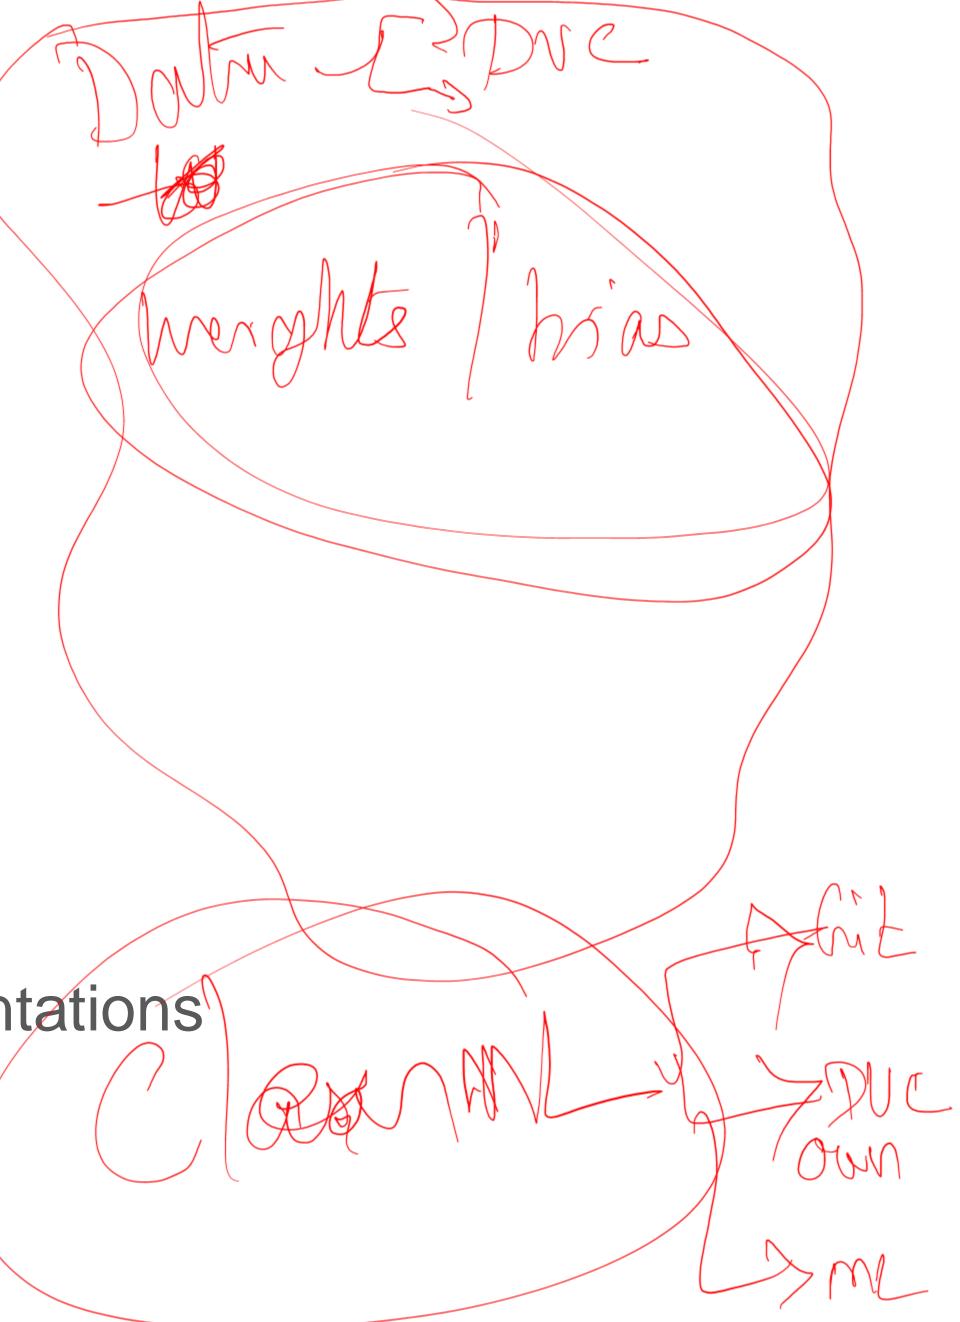
# MLOps Level 0: Manual Process

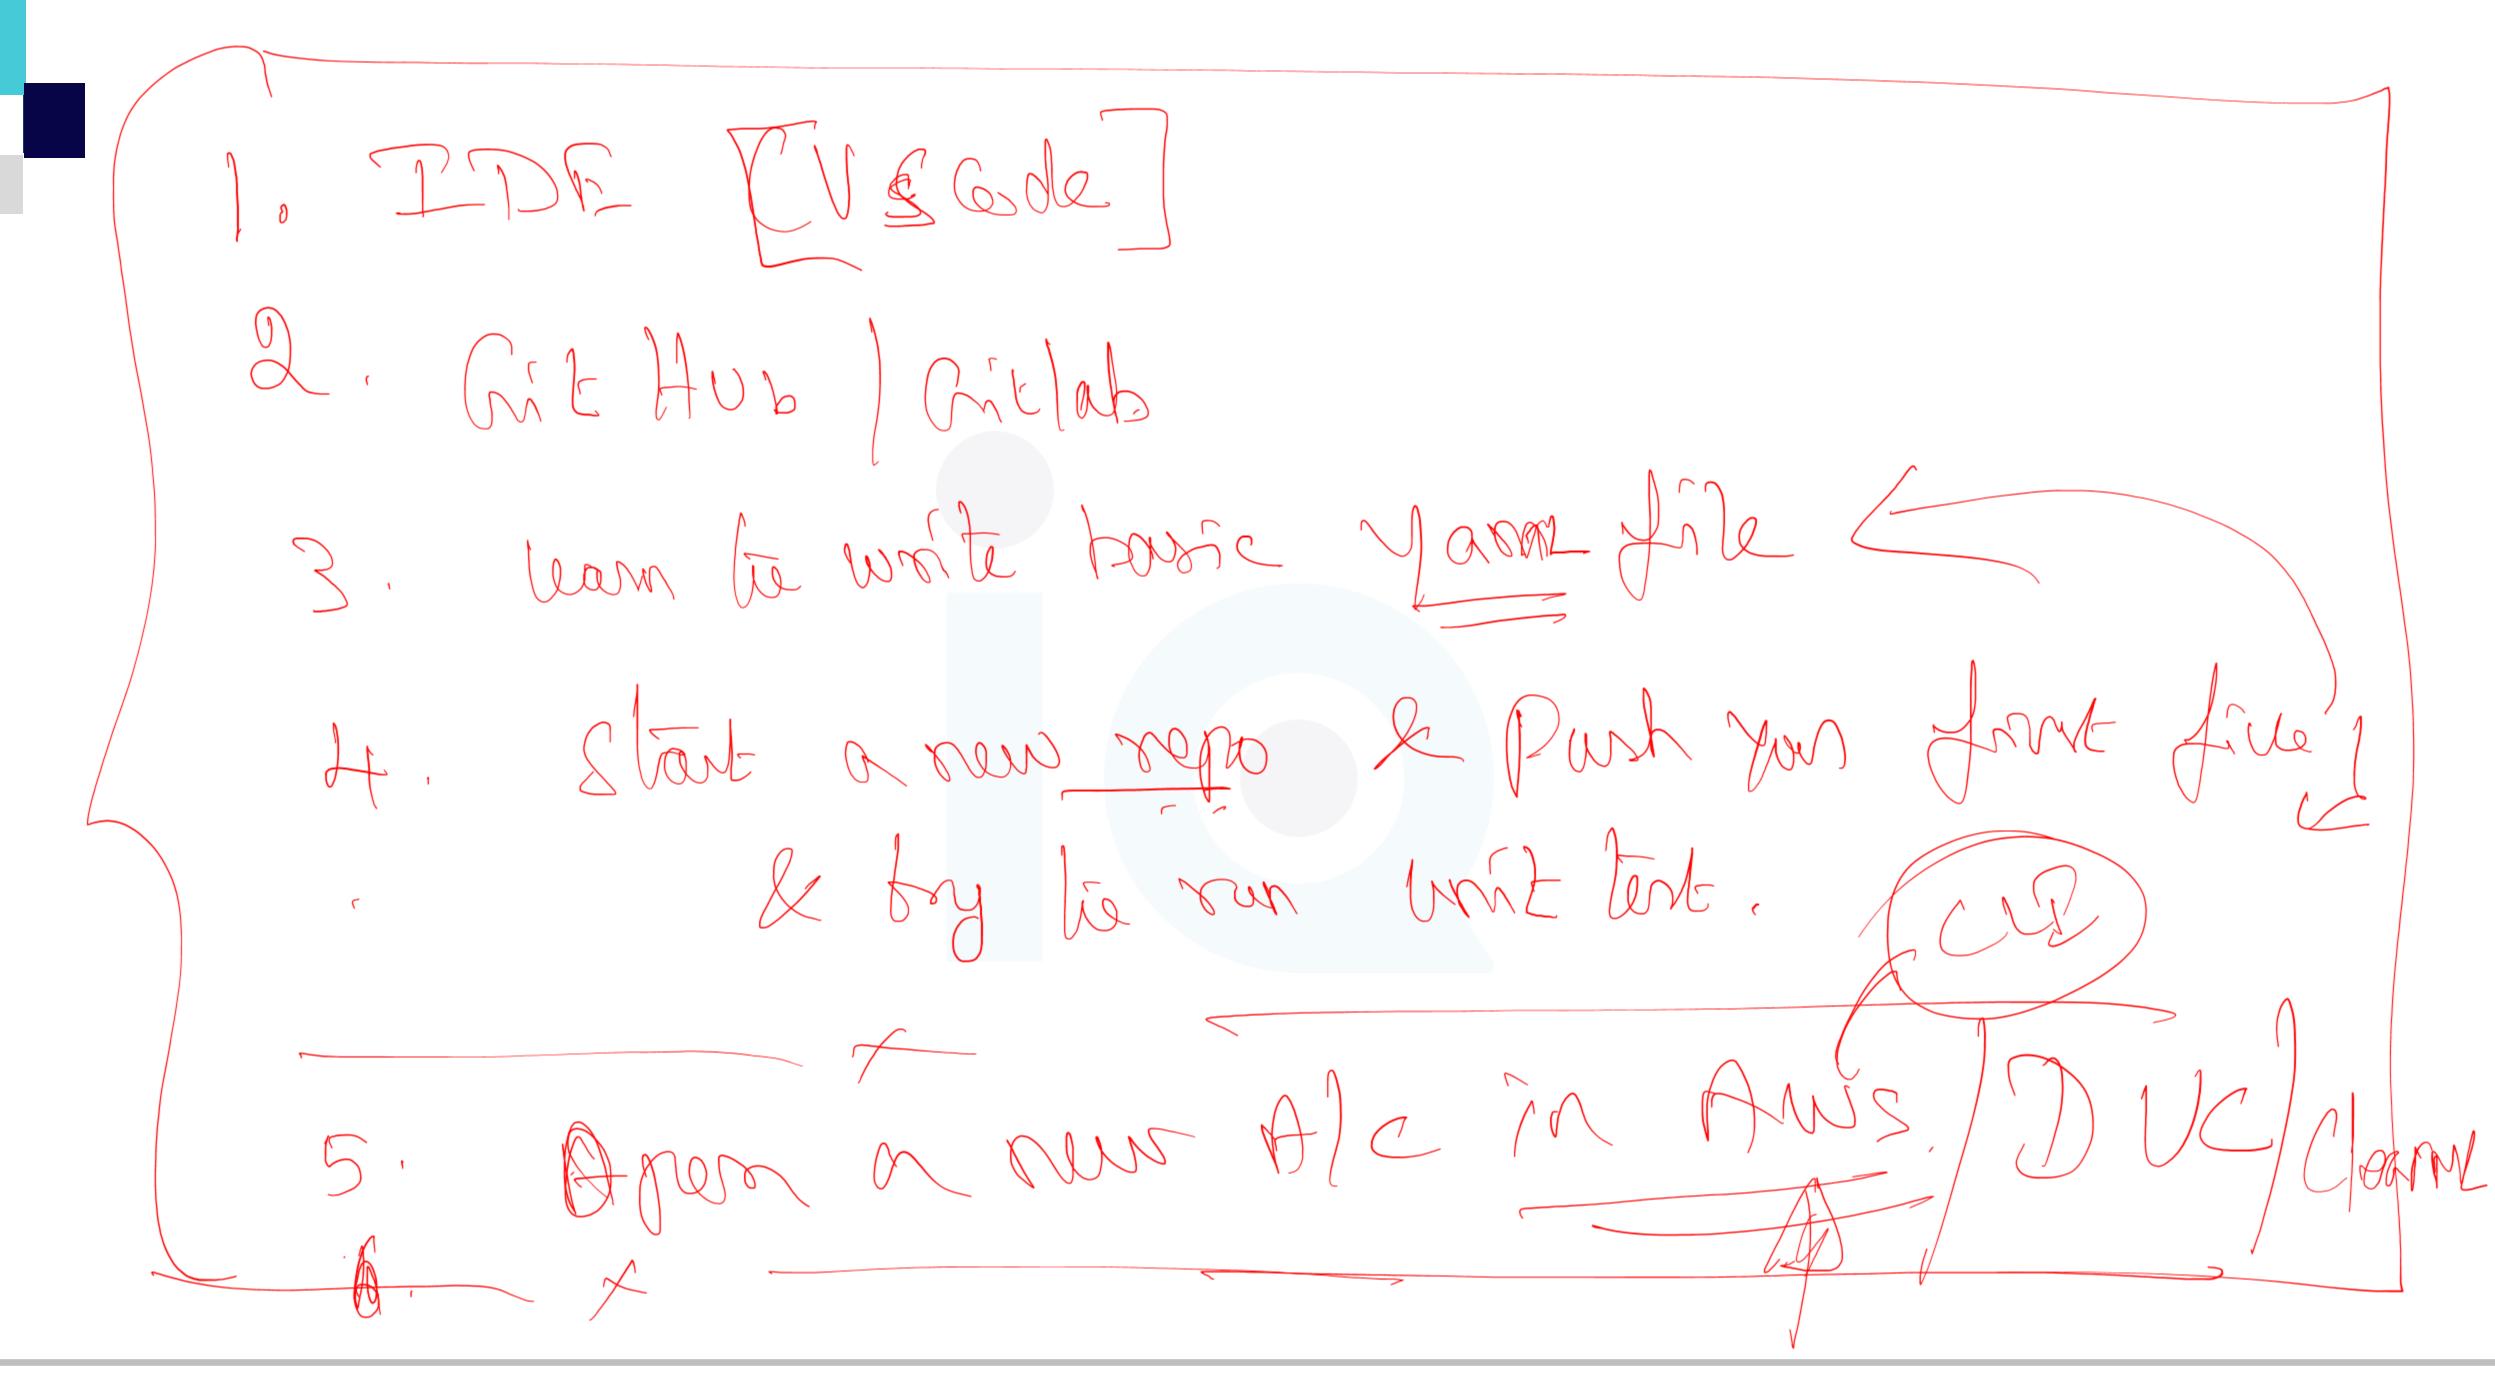
#### When to use?

- Common in many businesses
- Sufficient when the models change seldom

#### Recommendations

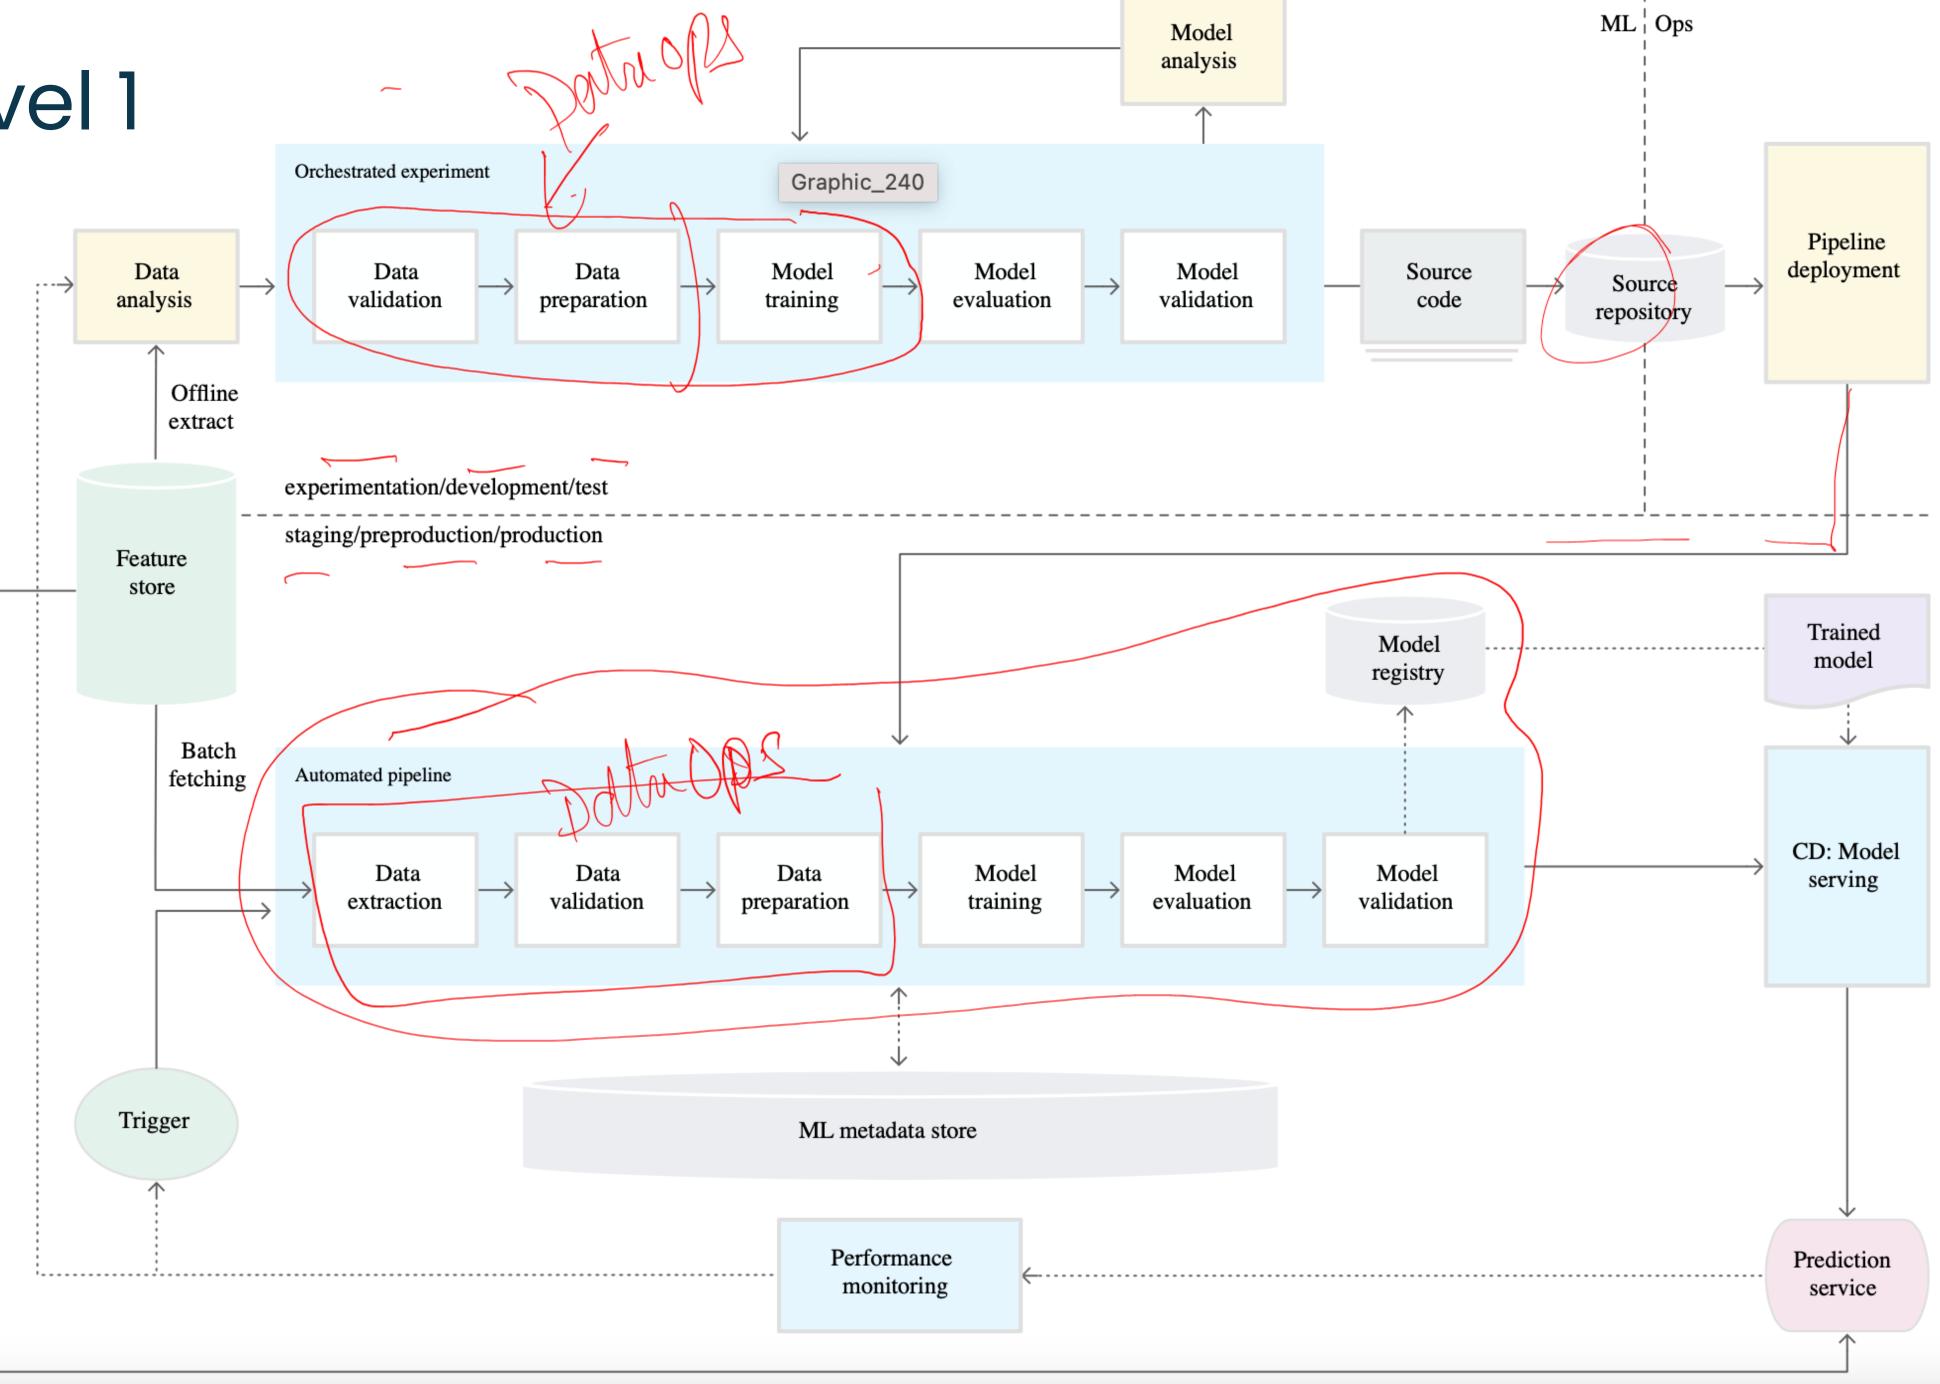
- Actively monitoring the quality of the model
- Frequently retraining the model, if needed
- Continuously experiment with new implementations
- Automate the process
  - Start implementing CI, CD, and CM/CT





MLOps Level 1

Level 1:
ML pipeline
automation



# MLOps Level 1: ML pipeline automation

- Perform continuous training of the model by automating the ML pipeline
- Retraining
  - Introduce automated data and model validation
  - Implement pipeline triggers and metadata management
- Rapid experiments
  - ML experiment is orchestrated
  - Better readiness to move to the whole pipeline to production
- CM/CT automation
  - Model in production is automatically retrained
  - Development and experiment environment is used in preproduction and production environment

# MLOps Level 1: ML pipeline automation

- CD of models
  - New models that are trained on new data are available
- Pipeline deployment
  - Unlike in Level 0, the whole training pipeline is deployed

# MLOps Level 1: Data and Model Validation

- Data Validation
  - Needed to decide whether to retrain the model or stop the execution of the pipeline
  - Data schema skews
    - Due to anomalies in the data
    - Data for downstream pipeline (data processing and model training) does not comply with expected schema
    - Stop the execution of pipeline, and let data scientist team to investigate the data
  - Data Value Skew: Significant changes in the statistical properties of data

# MLOps Level 1

#### Data and Model Validation

- Model Validation
  - Evaluate the retrained model before promoting to production
  - Produce evaluation metric using trained model on a test dataset
  - Compare the predictive metric of retrained model with current model (production model, baseline model and businessrequirement model)
  - Check the metric with various segment of the data
  - Test deployment-readiness, incl. infrastructure compatibility

# MLOps Level 1

#### **Feature Store**

 A centralized repository for standardized definitions, storage and access of features for training and serving

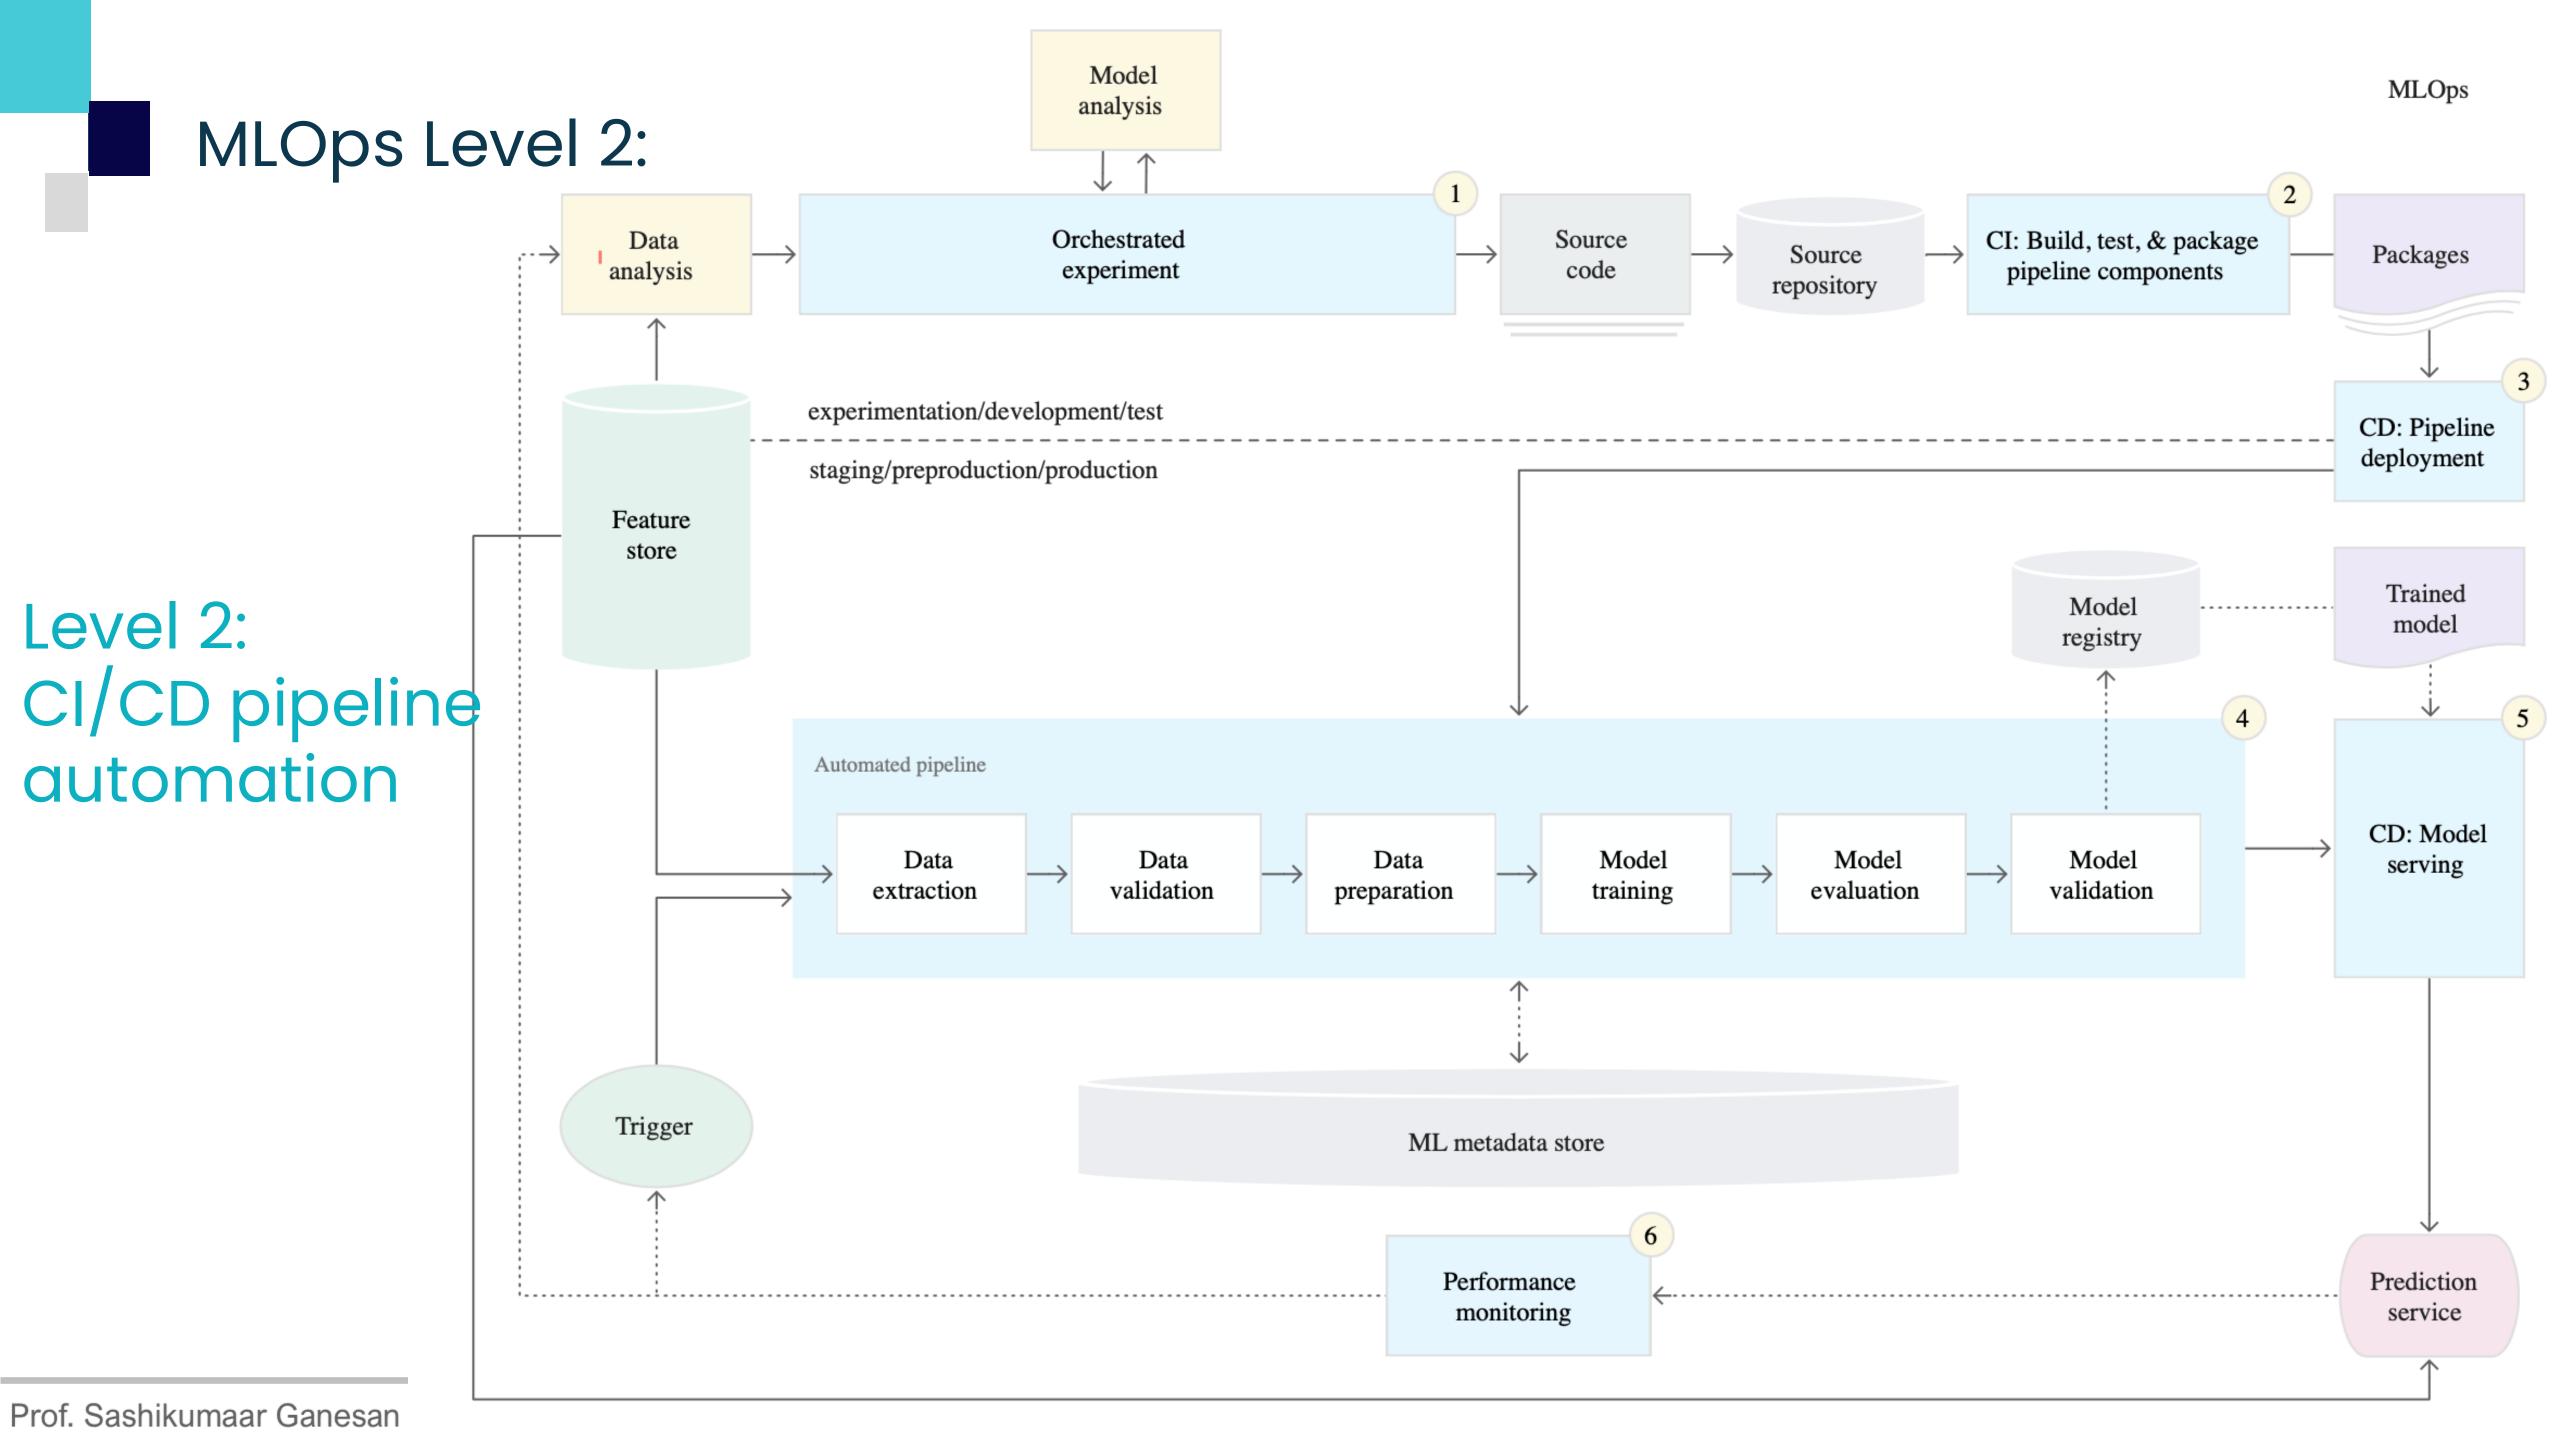
### Metadata management

- Store information about each execution
  - The pipeline and component versions used
  - Start and end time, date
  - Execution time of each step in the pipeline
  - Executor of pipeline
  - Parameter arguments that were passed to the pipeline, etc.

# MLOps Level 1:

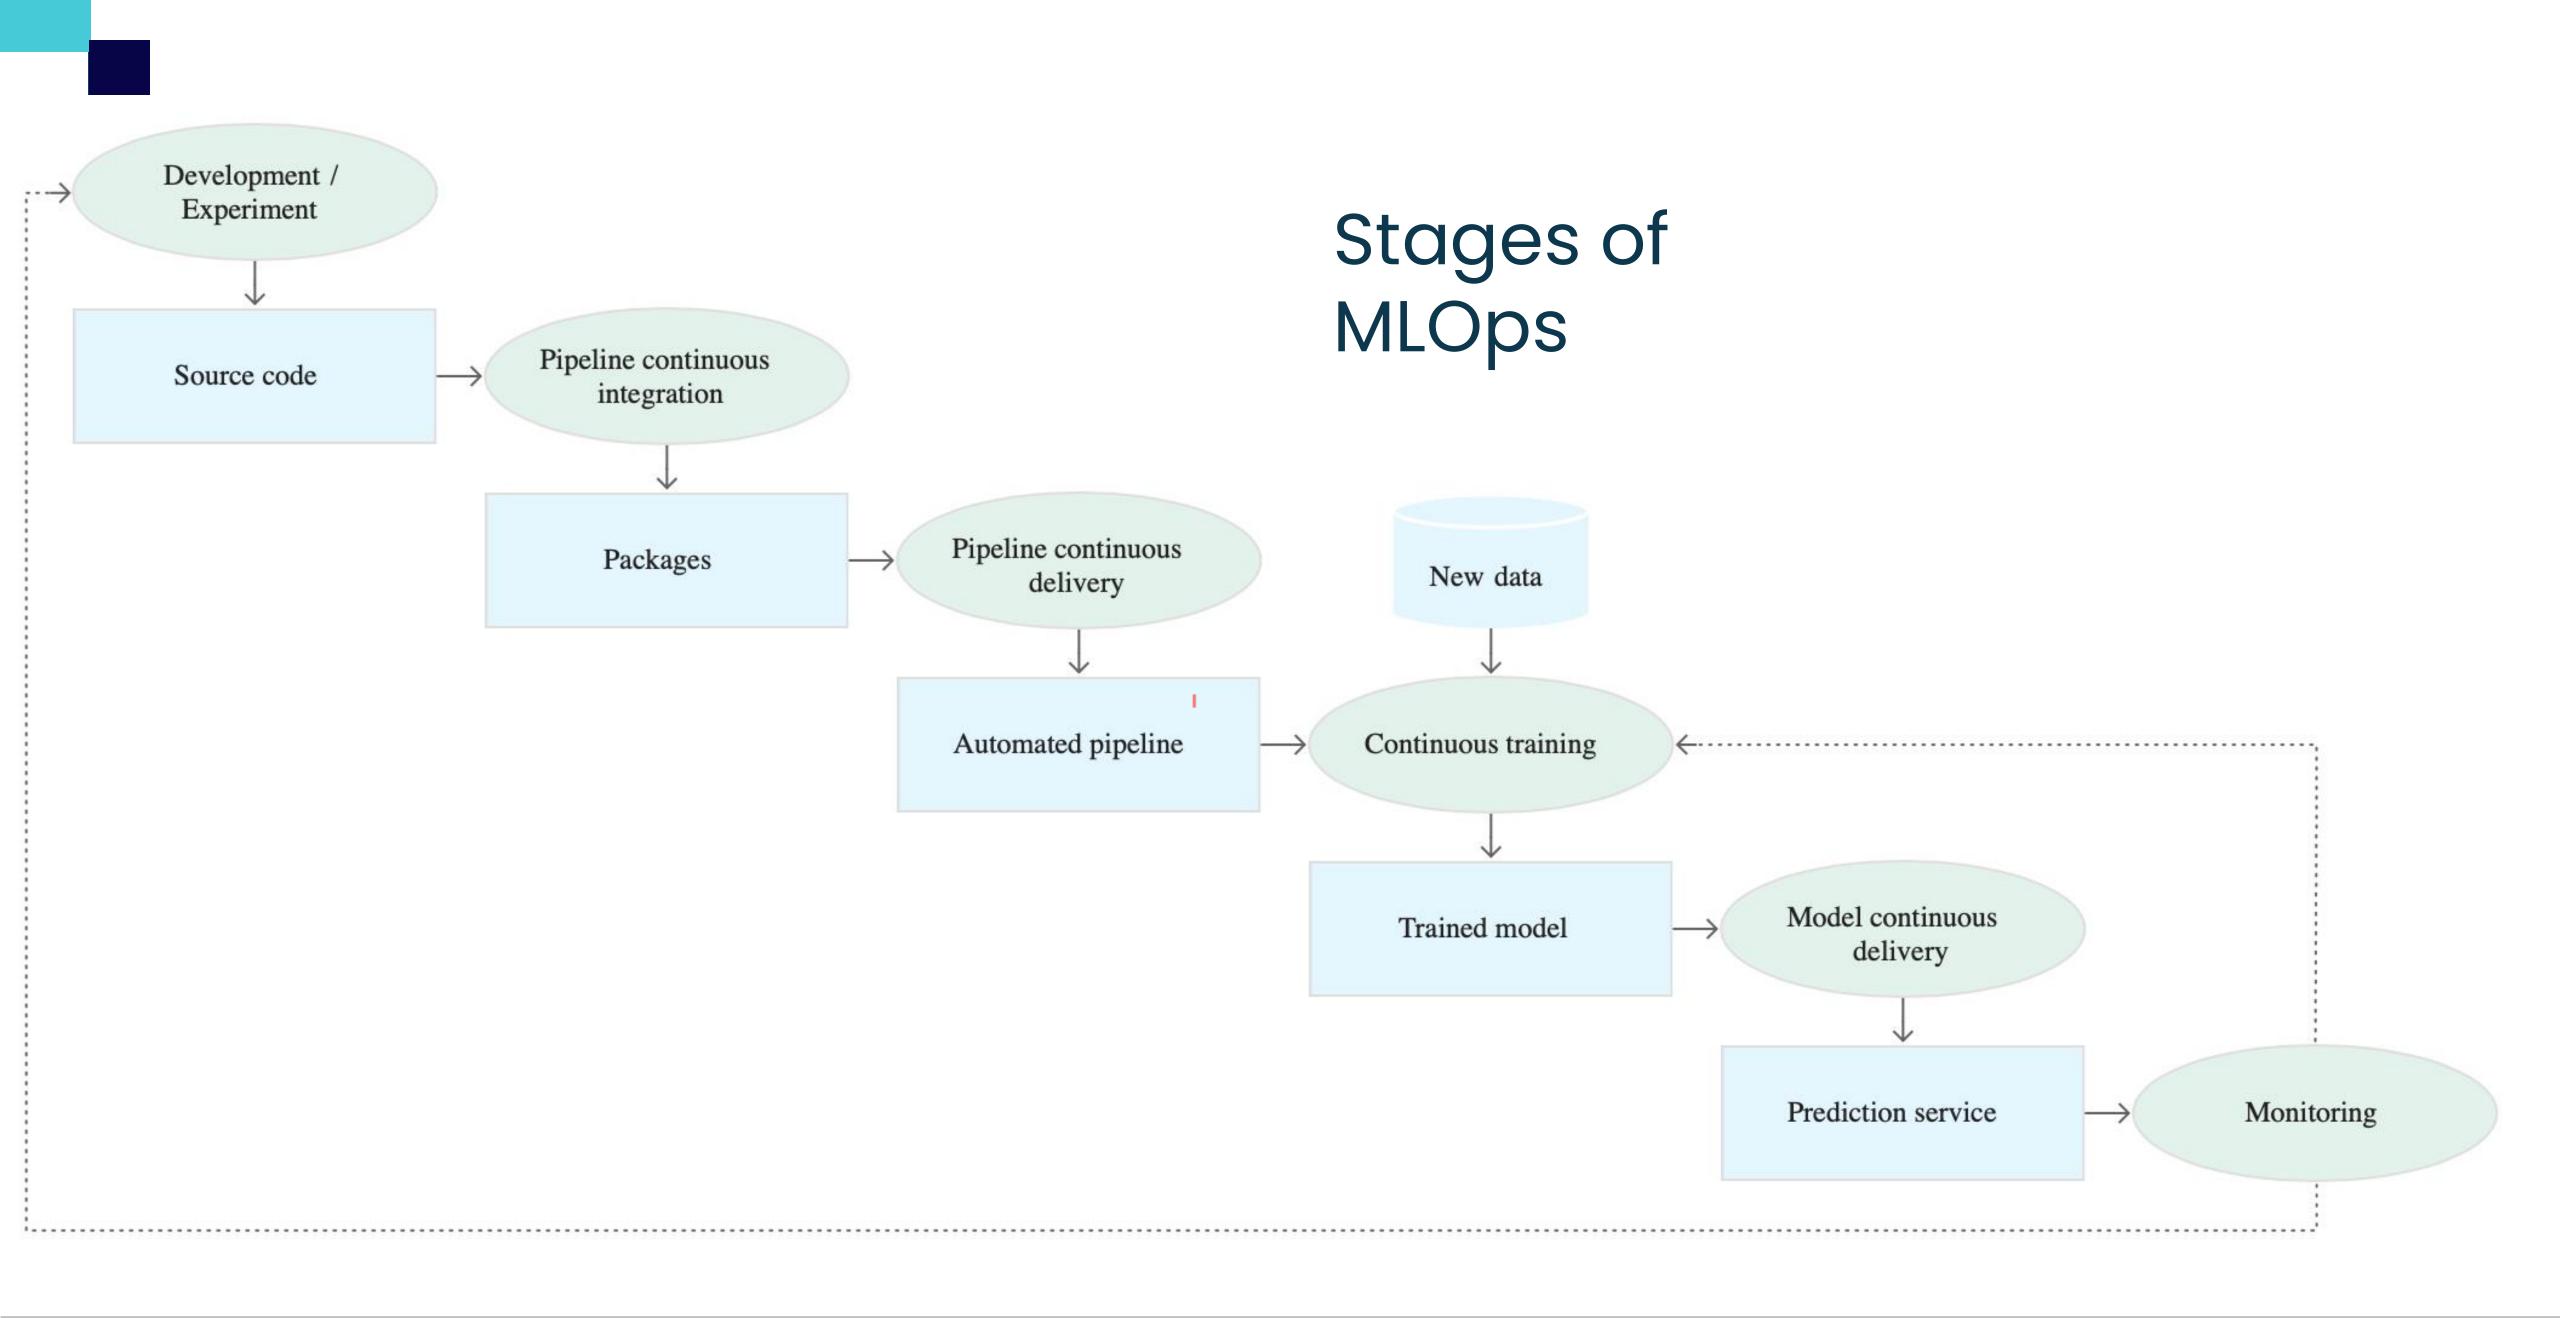
### ML pipeline triggers

- On demand: Ad-hoc manual execution of the pipeline
- On a schedule: Daily, weekly or monthly basis whenever new, labelled data available
- On availability of new data
- On model performance degradation
- On significant changes in the data distributions



# MLOps Level 2: CI/CD pipeline automation

- Allows to rapidly explore new ideas around feature engineering, model architecture and hyperparameters
- Automatically build, test and deploy new pipeline components to the target environments
- Steps include
  - Source control
  - Test and build services
  - Deployment services
  - Model registry
  - Feature store
  - ML metadata store
  - ML pipeline orchestrator



# MLOps Level 2: CI/CD pipeline automation

- Development and experimentation:
  - iteratively try new ML models where the experiment steps are orchestrated
  - source code is then pushed to a source repository
- Pipeline continuous integration:
  - run various tests
  - pipeline components (packages, executables, and artifacts) are deployed in a later stage
- Pipeline continuous delivery:
  - deploy the artifacts produced by the CI stage to the target environment

# MLOps Level 2: CI/CD pipeline automation

- Automated triggering:
  - Pipeline is automatically executed in production based on a schedule or in response to a trigger
  - Trained model is pushed to the model registry
- Model continuous delivery:
  - Serve the trained model as a prediction service for the predictions
  - The output of this stage is a deployed model prediction service
- Monitoring:
  - Collect statistics on the model performance based on live data.
  - Output of this stage is a trigger to execute the pipeline or to execute a new experiment cycle

# Summary

- Implementation and Automation
  - Continuous Integration (CI)
  - Continuous Delivery (CD)
  - Continuous Monitoring/Training (CM/CT)
  - MLOps Levels
    - Level 0: Manual Process
    - Level 1: ML Pipeline automation
    - Level 2: CI, CD and CT Pipeline automation

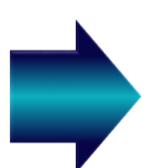
# Starting a new MLOps Project

• https://drivendata.github.io/cookiecutter-data-science/

## Building a MLOps Project: Phase 1: Planning & Design

# Define Objectives & Scope

- Clearly articulate project goals, target metrics (e.g., accuracy, latency, throughput), and desired business outcomes
- Identify stakeholders and their roles
- Create a project timeline with milestones



#### **Data Strategy**

- Data
  Sources: Determine data
  sources
  (internal, external, realtime, batch)
- Data Quality: Define data quality standards and establish cleansing/validation processes
- Data Labeling: Plan for data labeling if needed (inhouse, outsourced, crowdsourced)
- Data
   Governance: Establish
   policies for data



# Model Selection & Architecture Design

- Problem Type: Classify the ML problem (regression, classification, clustering, etc.)
- Algorithm
  Research: Explore
  potential algorithms and
  architectures based on the
  problem type and data
  characteristics
- Baseline Model: Create a simple model to establish a performance baseline
- Model Architecture: Design a

access, security, and

# Building a MLOps Project: Phase 2: Development & Experimentation

### DataOps Implementation

- Data Pipelines: Create automated pipelines for data ingestion, cleaning, transformation, and feature engineering
- Data Versioning: Use tools like DVC or Pachyderm to track changes to data and pipelines
- Feature Store:

   (Optional) Implement a feature store like Feast or
   Tecton to centralize and manage features
- Data Exploration and
   Visualization: Explore



- Experiment Tracking:
  Log experiments with
  MLflow, TensorBoard, or
  Weights & Biases to track
  model hyperparameters,
  metrics, and artifacts
- Model Versioning: Use tools like MLflow to version models and track performance over time
- Hyperparameter
  Tuning: Perform
  systematic tuning of model
  hype
  optim

  CLEARIML
  optim



# Version Control and Collaboration

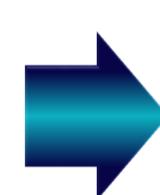
- Git: Use Git for code versioning and collaboration
- Code Reviews: Implement code reviews for quality assurance
- Branching Strategy:

   Adopt a branching strategy
   (e.g., Gitflow) to manage
   code changes effectively
- Data and Model
  Testing: In MLOps, it's crucial to test not only your code but also your data and model performance

## Building a MLOps Project: Phase 3: Deployment and Monitoring

#### Containerization

- **Docker:** Containerize the model and its dependencies using Docker
- Docker Registries: Store container images in a registry like Docker Hub or a private registry
- Kubernetes (Module 7):
  - Cluster Setup: Set up a Kubernetes cluster (on-premises or in the cloud).
  - Deployment Manifests: Create Kubernetes manifests (YAML files) to define Pods, Services, Deployments, and other resources for your ML application.
  - Scalability: Utilize Kubernetes' autoscaling capabilities to adjust resources based on demand.
  - Rollouts and Rollbacks: Use Kubernetes to manage gradual rollouts and rollbacks of model versions.



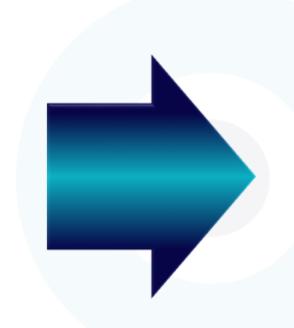
### CI/CD Pipeline

- Automated Testing: Write unit, integration, and regression tests
- CI/CD Infrastructure: Set up infrastructure for automated testing and deployment
- CI/CD Tools: Use tools like GitHub Actions, GitLab CI, Jenkins, or CircleCI

## Building a MLOps Project: Phase 3: Deployment and Monitoring

### Deployment

- Cloud Deployment: Deploy to cloud platforms like AWS SageMaker, Azure ML, or Google Cloud AI Platform
- Model Serving: Use tools like KFServing,
   Seldon Core, or BentoML for model serving
- Web Applications: Create web applications (e.g., with Streamlit, Gradio, Flask) for model interaction
- Edge Deployment: (Optional) Deploy models to edge devices if applicable



#### Continuous Monitoring and Maintenance

- Logging: Implement logging of model inputs, outputs, and performance metrics
- Monitoring Tools: Use tools like Prometheus, Grafana, or cloud-specific monitoring solutions
- Alerting: Set up alerts for performance degradation or anomalies
- Model Retraining: Establish a process for retraining models with fresh data

## Building a MLOps Project: Phase 4: Optimization and Governance

### Model Optimization

- Performance
   Optimization:
   Profile the model
   to identify
   bottlenecks and
   optimize for speed
   and resource
   usage
- Model
   Compression:
   Apply techniques
   like quantization or
   pruning to reduce
   model size
- Batch
   Prediction: Use

# Security and Governance

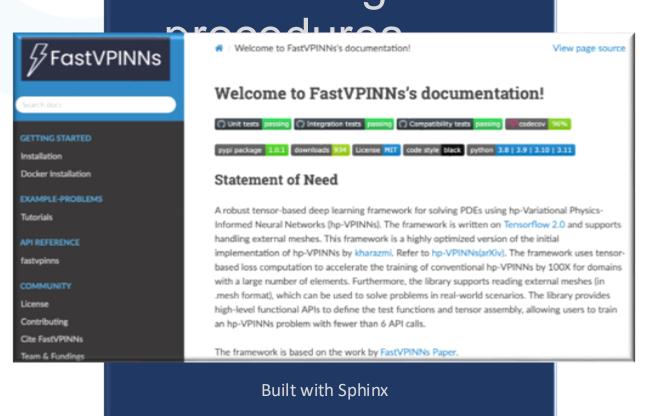
- Model Security:
  Implement security
  measures to
  protect the model
  from unauthorized
  access and
  adversarial attacks
- Model Bias and Fairness:
   Regularly audit the model for bias and take corrective

   measures
- Model
   Explainability:
   Use tools like

CHAD OF HIME to

#### Documentation

 Create thorough documentation for the project, including data descriptions, model architectures, deployment instructions, and monitoring



# Feedback & Iteration

- Collect
   Feedback:
   Gather feedback
   from users and
   stakeholders
- Analyze
   Performance:
   Continuously analyze model performance metrics and user feedback
- Iterate: Refine the model, data pipelines, and deployment process based on feedback and analysis





(Hands-on session in the afternoon)

- What is Docker?
  - A platform for developing, shipping, and running applications.
  - Utilizes containerization to make applications portable and consistent across different environments.
  - Containers encapsulate an application with all of its dependencies.

- Relevance in Modern Software Development
  - Facilitates continuous integration and continuous deployment (CI/CD) by ensuring that software runs the same in all environments.
  - Reduces "it works on my machine" problems by providing a consistent environment from development to production.
  - Enables microservices architecture by allowing each service to be containerized and scaled independently.
  - Streamlines development by allowing developers to create predictable and efficient work environments.
  - Enhances collaboration between development and operations teams for faster and more reliable software delivery.

## Docker: Containers vs. Virtual Machines

### Containers

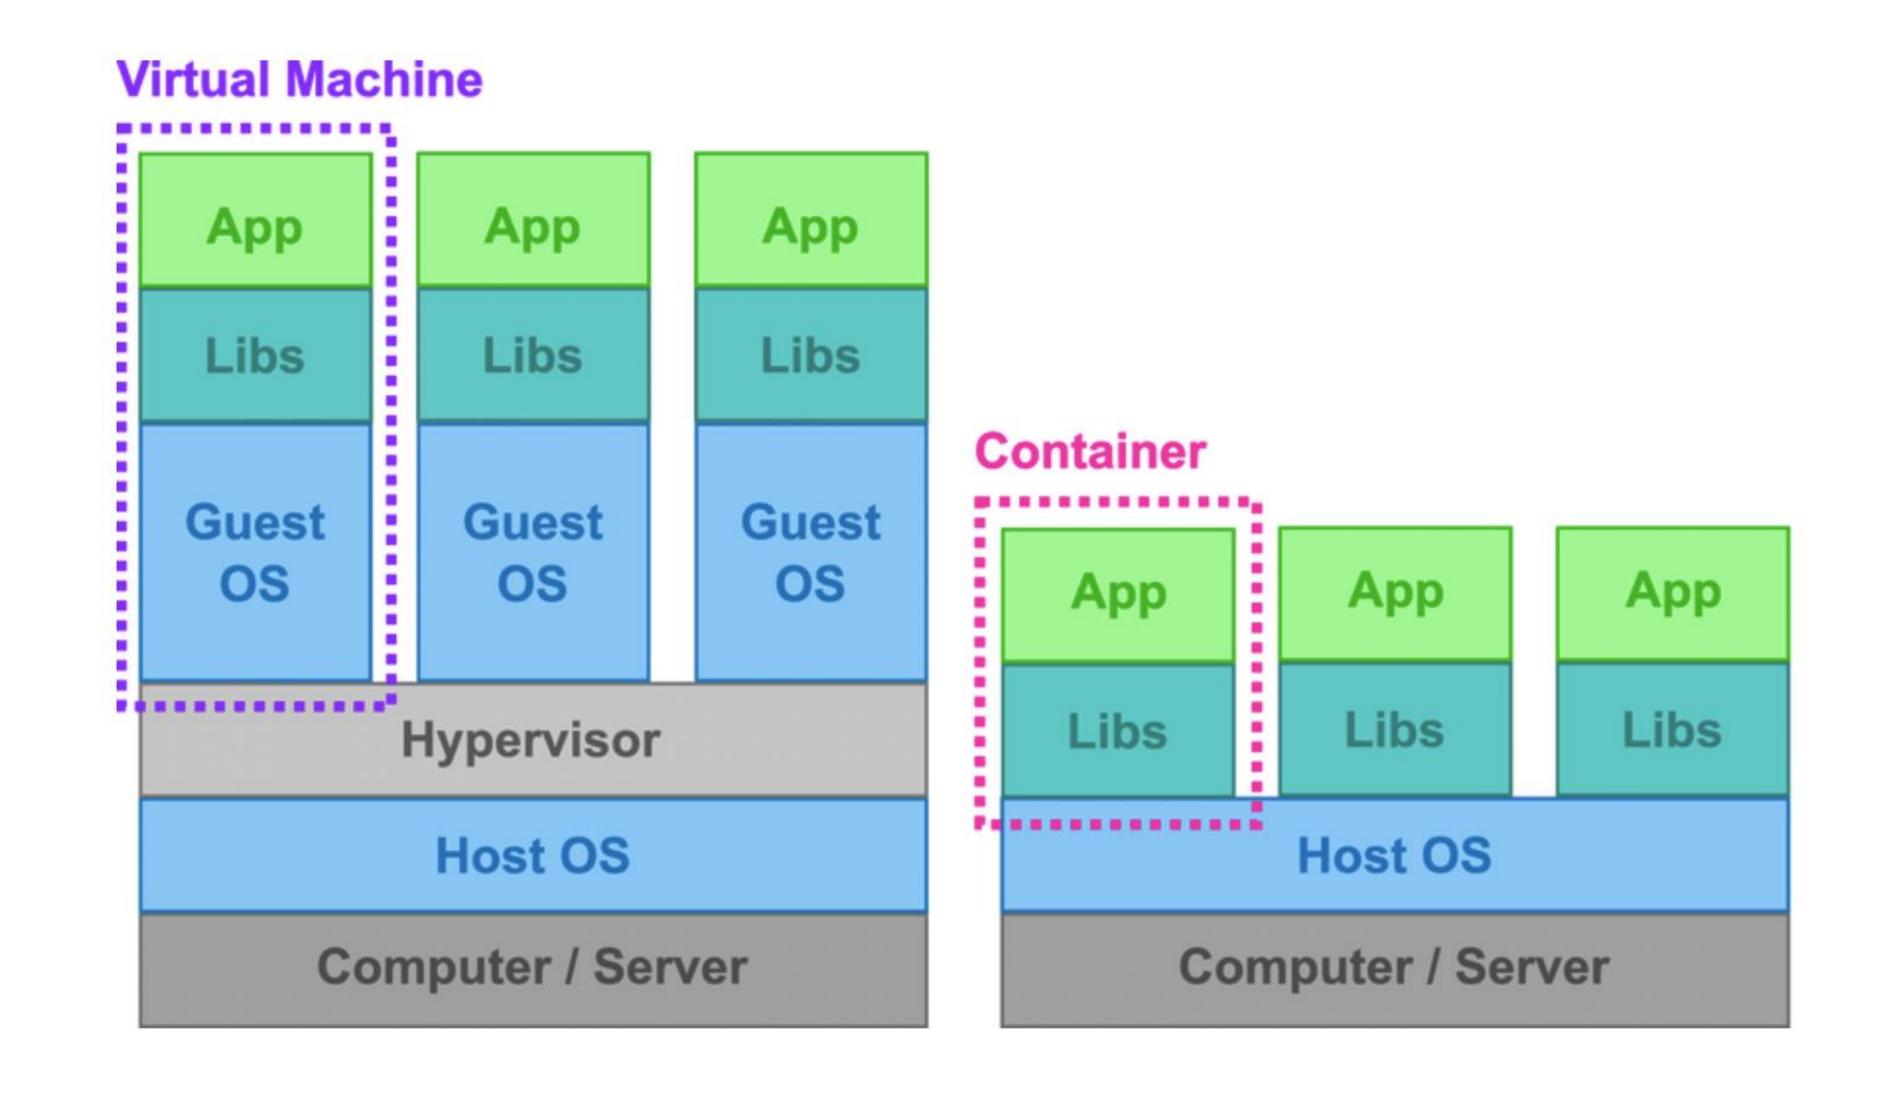
- Lightweight: Containers are much smaller and more resource-efficient than VMs because they share the underlying host operating system (OS) kernel.
- Portable: Containers package applications with all their dependencies, ensuring consistent behavior across different environments.
- Fast Startup: Since containers don't need to boot up a full OS, they start up almost instantly.
- Efficient Resource Usage: Containers share the host OS resources, reducing overhead and allowing more applications to run on the same hardware.
- Suitable for Microservices: Containers are well-suited for breaking down applications into smaller, independent components (microservices), which

## Docker: Containers vs. Virtual Machines

### Virtual Machines

- Full OS: Each VM runs its own complete guest operating system, providing strong isolation between VMs.
- More Overhead: VMs require more resources than containers due to the need to run a full OS for each VM.
- Slower Startup: Booting up a VM takes longer than starting a container.
- Stronger Isolation: VMs provide better security and fault isolation than containers due to the separation provided by the hypervisor.
- Suitable for Legacy Applications: VMs are often used to run legacy applications that may not be easily containerized.

## Docker: Containers vs. Virtual Machines



# Docker: Core Concepts

- Images and Containers
  - Docker images are lightweight, stand-alone, executable software packages that include everything needed to run a piece of software, including the code, runtime, libraries, environment variables, and config files.
  - Containers are a runtime instance of Docker images an image becomes a container when it runs on Docker Engine.
  - Containers are isolated from each other and the host system, but can communicate through well-defined channels.

# Docker: Core Concepts

### Docker Hub

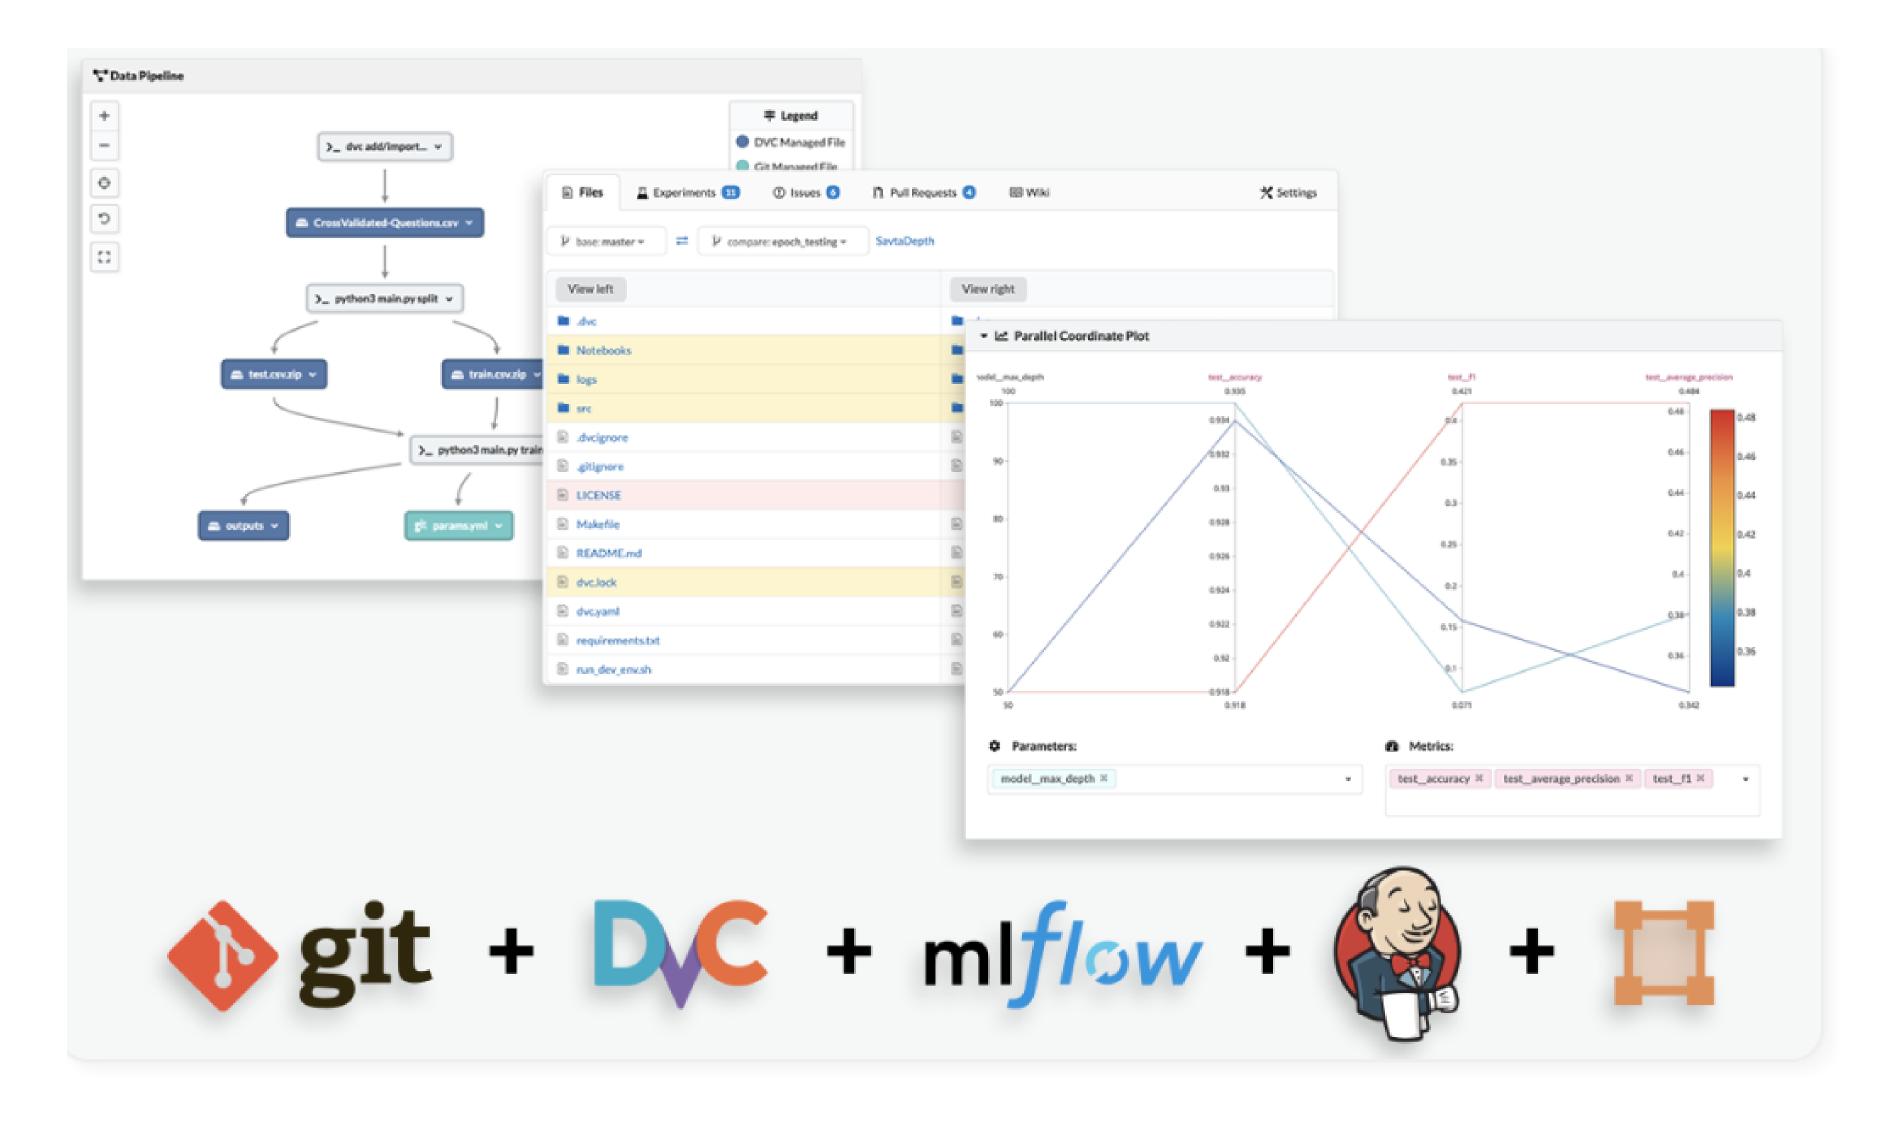
- Docker Hub is a cloud-based registry service that allows you to link to code repositories, build your images, test them, store manually pushed images, and link to Docker Cloud.
- Provides a comprehensive repository for Docker container images with both public and private storage options.
- It is the default registry where Docker looks for images.
- Offers automated build capabilities for creating Docker images from online source code repositories.

- Docker's Benefits
  - Promotes consistency across development, staging, and production environments.
  - Ensures that applications run the same, regardless of where they are deployed.
  - Reduces overhead and increases efficiency compared to traditional virtual machines.

- Docker in the Development Toolchain
  - Integrates with GitHub for version control, enabling seamless code updates and tracking.
  - Works alongside GitHub Actions for automated testing and deployment pipelines.
  - Complements AWS Compute services to provide scalable and secure hosting for containerized applications.
  - Facilitates data version control with DVC, making it easier to track and manage datasets and machine learning models.
  - Enhances ML operations by working with MLflow to manage the lifecycle of machine learning models, including experimentation, reproducibility, and deployment.

Prof. Sashikumaar Ganesan

- Synergy Between Technologies
  - Containerization with Docker offers a standardized unit for software development, enabling a smooth workflow across tools like GitHub, AWS, DVC, and MLflow.
  - Simplifies the complexity of managing dependencies and environments in machine learning projects.
  - Empowers teams to build, test, and release software faster and more reliably, contributing to the DevOps culture of collaboration and efficiency.



Prof. Sashikumaar Ganesan IISc, Bangalore | Zenteiq Aitech Innovations